

A Machine Learning Approach to Coreference Resolution of Noun Phrases

Wee Meng Soon*
DSO National Laboratories

Hwee Tou Ng†
DSO National Laboratories

Daniel Chung Yong Lim‡
DSO National Laboratories

In this paper, we present a learning approach to coreference resolution of noun phrases in unrestricted text. The approach learns from a small, annotated corpus and the task includes resolving not just a certain type of noun phrase (e.g., pronouns) but rather general noun phrases. It also does not restrict the entity types of the noun phrases; that is, coreference is assigned whether they are of “organization,” “person,” or other types. We evaluate our approach on common data sets (namely, the MUC-6 and MUC-7 coreference corpora) and obtain encouraging results, indicating that on the general noun phrase coreference task, the learning approach holds promise and achieves accuracy comparable to that of nonlearning approaches. Our system is the first learning-based system that offers performance comparable to that of state-of-the-art nonlearning systems on these data sets.

1. Introduction

Coreference resolution is the process of determining whether two expressions in natural language refer to the same entity in the world. It is an important subtask in natural language processing systems. In particular, information extraction (IE) systems like those built in the DARPA Message Understanding Conferences (Chinchor 1998; Sundheim 1995) have revealed that coreference resolution is such a critical component of IE systems that a separate coreference subtask has been defined and evaluated since MUC-6 (MUC-6 1995).

In this paper, we focus on the task of determining coreference relations as defined in MUC-6 (MUC-6 1995) and MUC-7 (MUC-7 1997). Specifically, a coreference relation denotes an identity of reference and holds between two textual elements known as markables, which can be definite noun phrases, demonstrative noun phrases, proper names, appositives, sub-noun phrases that act as modifiers, pronouns, and so on. Thus, our coreference task resolves general noun phrases and is not restricted to a certain type of noun phrase such as pronouns. Also, we do not place any restriction on the possible candidate markables; that is, all markables, whether they are “organization,” “person,” or other entity types, are considered. The ability to link coreferring noun phrases both within and across sentences is critical to discourse analysis and language understanding in general.

* 20 Science Park Drive, Singapore 118230. E-mail: sweemeng@dso.org.sg

† 20 Science Park Drive, Singapore 118230. E-mail: nhweetou@dso.org.sg. This author is also affiliated with the Department of Computer Science, School of Computing, National University of Singapore. Web address: <http://www.comp.nus.edu.sg/~nght>

‡ 20 Science Park Drive, Singapore 118230. E-mail: ichungyo@dso.org.sg

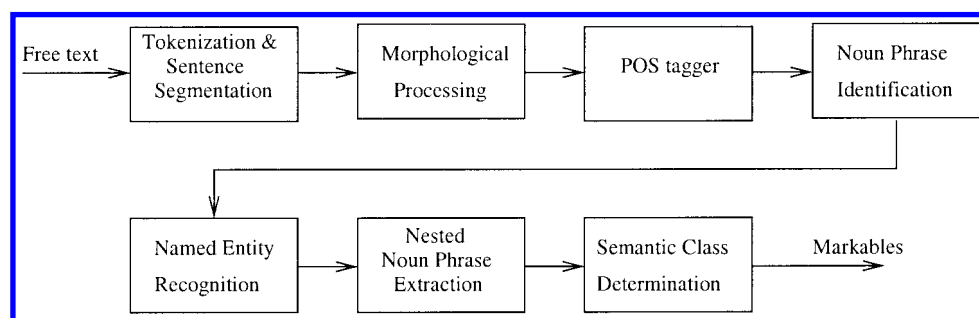


Figure 1
System architecture of natural language processing pipeline.

2. A Machine Learning Approach to Coreference Resolution

We adopt a corpus-based, machine learning approach to noun phrase coreference resolution. This approach requires a relatively small corpus of training documents that have been annotated with coreference chains of noun phrases. All possible markables in a training document are determined by a pipeline of language-processing modules, and training examples in the form of feature vectors are generated for appropriate pairs of markables. These training examples are then given to a learning algorithm to build a classifier. To determine the coreference chains in a new document, all markables are determined and potential pairs of coreferring markables are presented to the classifier, which decides whether the two markables actually corefer. We give the details of these steps in the following subsections.

2.1 Determination of Markables

A prerequisite for coreference resolution is to obtain most, if not all, of the possible markables in a raw input text. To determine the markables, a pipeline of natural language processing (NLP) modules is used, as shown in Figure 1. They consist of tokenization, sentence segmentation, morphological processing, part-of-speech tagging, noun phrase identification, named entity recognition, nested noun phrase extraction, and semantic class determination. **As far as coreference resolution is concerned, the goal of these NLP modules is to determine the boundary of the markables, and to provide the necessary information about each markable for subsequent generation of features in the training examples.**

Our part-of-speech tagger is a standard statistical tagger based on the Hidden Markov Model (HMM) (Church 1988). Similarly, we built a statistical HMM-based noun phrase identification module that determines the noun phrase boundaries solely based on the part-of-speech tags assigned to the words in a sentence. We also implemented a module that recognizes MUC-style named entities, that is, organization, person, location, date, time, money, and percent. Our named entity recognition module uses the HMM approach of Bikel, Schwartz, and Weischedel (1999), which learns from a tagged corpus of named entities. That is, our part-of-speech tagger, noun phrase identification module, and named entity recognition module are all based on HMMs and learn from corpora tagged with parts of speech, noun phrases, and named entities, respectively. Next, both the noun phrases determined by the noun phrase identification module and the named entities are merged in such a way that if the noun phrase overlaps with a named entity, the noun phrase boundaries will be adjusted to subsume the named entity.

The nested noun phrase extraction module subsequently accepts the noun phrases and determines the nested phrases for each noun phrase. The nested noun phrases are divided into two groups:

1. Nested noun phrases from possessive noun phrases. Consider two possessive noun phrases marked by the noun phrase module, *his long-range strategy* and *Eastern's parent*. The nested noun phrase for the first phrase is the pronoun *his*, while for the second one, it is the proper name *Eastern*.
2. Nested noun phrases that are modifier nouns (or prenominals). For example, the nested noun phrase for *wage reductions* is *wage*, and for *Union representatives*, it is *Union*.

Finally, the markables needed for coreference resolution are the union of the noun phrases, named entities, and nested noun phrases found. For markables without any named entity type, semantic class is determined by the semantic class determination module. More details regarding this module are given in the description of the semantic class agreement feature.

To achieve acceptable recall for coreference resolution, it is most critical that the eligible candidates for coreference be identified correctly in the first place. In order to test our system's effectiveness in determining the markables, we attempted to match the markables generated by our system against those appearing in the coreference chains annotated in 100 SGML documents, a subset of the training documents available in MUC-6. We found that our system is able to correctly identify about 85% of the noun phrases appearing in coreference chains in the 100 annotated SGML documents. Most of the unmatched noun phrases are of the following types:

1. Our system generated a head noun that is a subset of the noun phrase in the annotated corpus. For example, *Saudi Arabia, the cartel's biggest producer* was annotated as a markable, but our system generated only *Saudi Arabia*.
2. Our system extracted a sequence of words that cannot be considered as a markable.
3. Our system extracted markables that appear to be correct but do not match what was annotated. For example, our system identified *selective wage reductions*, but *wage reductions* was annotated instead.

2.2 Determination of Feature Vectors

To build a learning-based coreference engine, we need to devise a set of features that is useful in determining whether two markables corefer or not. In addition, these features must be generic enough to be used across different domains. Since the MUC-6 and MUC-7 tasks define coreference guidelines for all types of noun phrases and different types of noun phrases behave differently in terms of how they corefer, our features must be able to handle this and give different coreference decisions based on different types of noun phrases. In general, there must be some features that indicate the type of a noun phrase. Altogether, we have five features that indicate whether the markables are definite noun phrases, demonstrative noun phrases, pronouns, or proper names.

There are many important knowledge sources useful for coreference. We wanted to use those that are not too difficult to compute. One important factor is the distance

between the two markables. McEnery, Tanaka, and Botley (1997) have done a study on how distance affects coreference, particularly for pronouns. One of their conclusions is that the antecedents of pronouns do exhibit clear quantitative patterns of distribution.

The distance feature has different effects on different noun phrases. For proper names, locality of the antecedents may not be so important. We include the distance feature so that the learning algorithm can best decide the distribution for different classes of noun phrases.

There are other features that are related to the gender, number, and semantic class of the two markables. Such knowledge sources are commonly used for the task of determining coreference.

Our feature vector consists of a total of 12 features described below, and is derived based on two extracted markables, *i* and *j*, where *i* is the potential antecedent and *j* is the anaphor. Information needed to derive the feature vectors is provided by the pipeline of language-processing modules prior to the coreference engine.

1. **Distance Feature (DIST):** Its possible values are 0, 1, 2, 3, This feature captures the distance between *i* and *j*. If *i* and *j* are in the same sentence, the value is 0; if they are one sentence apart, the value is 1; and so on.
2. **i-Pronoun Feature (I_PRONOUN):** Its possible values are true or false. If *i* is a pronoun, return true; else return false. Pronouns include reflexive pronouns (*himself, herself*), personal pronouns (*he, him, you*), and possessive pronouns (*hers, her*).
3. **j-Pronoun Feature (J_PRONOUN):** Its possible values are true or false. If *j* is a pronoun (as described above), then return true; else return false.
4. **String Match Feature (STR_MATCH):** Its possible values are true or false. If the string of *i* matches the string of *j*, return true; else return false. We first remove articles (*a, an, the*) and demonstrative pronouns (*this, these, that, those*) from the strings before performing the string comparison. Therefore, *the license* matches *this license*, *that computer* matches *computer*.
5. **Definite Noun Phrase Feature (DEF_NP):** Its possible values are true or false. In our definition, a definite noun phrase is a noun phrase that starts with the word *the*. For example, *the car* is a definite noun phrase. If *j* is a definite noun phrase, return true; else return false.
6. **Demonstrative Noun Phrase Feature (DEM_NP):** Its possible values are true or false. A demonstrative noun phrase is one that starts with the word *this, that, these, or those*. If *j* is a demonstrative noun phrase, then return true; else return false.
7. **Number Agreement Feature (NUMBER):** Its possible values are true or false. If *i* and *j* agree in number (i.e., they are both singular or both plural), the value is true; otherwise false. Pronouns such as *they* and *them* are plural, while *it, him*, and so on, are singular. The morphological root of a noun is used to determine whether it is singular or plural if the noun is not a pronoun.
8. **Semantic Class Agreement Feature (SEMCLASS):** Its possible values are true, false, or unknown. In our system, we defined the following semantic classes: "female," "male," "person," "organization," "location," "date," "time," "money," "percent," and "object." These semantic classes

are arranged in a simple ISA hierarchy. Each of the “female” and “male” semantic classes is a subclass of the semantic class “person,” while each of the semantic classes “organization,” “location,” “date,” “time,” “money,” and “percent” is a subclass of the semantic class “object.” Each of these defined semantic classes is then mapped to a WordNet synset (Miller 1990). For example, “male” is mapped to the second sense of the noun *male* in WordNet, “location” is mapped to the first sense of the noun *location*, and so on.

The semantic class determination module assumes that the semantic class for every markable extracted is the first sense of the head noun of the markable. Since WordNet orders the senses of a noun by their frequency, this is equivalent to choosing the most frequent sense as the semantic class for each noun. If the selected semantic class of a markable is a subclass of one of our defined semantic classes *C*, then the semantic class of the markable is *C*; else its semantic class is “unknown.”

The semantic classes of markables *i* and *j* are in agreement if one is the parent of the other (e.g., *chairman* with semantic class “person” and *Mr. Lim* with semantic class “male”), or they are the same (e.g., *Mr. Lim* and *he*, both of semantic class “male”). The value returned for such cases is true. If the semantic classes of *i* and *j* are not the same (e.g., *IBM* with semantic class “organization” and *Mr. Lim* with semantic class “male”), return false. If either semantic class is “unknown,” then the head noun strings of both markables are compared. If they are the same, return true; else return unknown.

9. **Gender Agreement Feature (GENDER):** Its possible values are true, false, or unknown. The gender of a markable is determined in several ways. Designators and pronouns such as *Mr.*, *Mrs.*, *she*, and *he*, can determine the gender. For a markable that is a person’s name, such as *Peter H. Diller*, the gender cannot be determined by the above method. In our system, the gender of such a markable can be determined if markables are found later in the document that refer to *Peter H. Diller* by using the designator form of the name, such as *Mr. Diller*. If the designator form of the name is not present, the system will look through its database of common human first names to determine the gender of that markable. The gender of a markable will be unknown for noun phrases such as *the president* and *chief executive officer*. The gender of other markables that are not “person” is determined by their semantic classes. Unknown semantic classes will have unknown gender while those that are objects will have “neutral” gender. If the gender of either markable *i* or *j* is unknown, then the gender agreement feature value is unknown; else if *i* and *j* agree in gender, then the feature value is true; otherwise its value is false.
10. **Both-Proper-Names Feature (PROPER_NAME):** Its possible values are true or false. A proper name is determined based on capitalization. Prepositions appearing in the name such as *of* and *and* need not be in uppercase. If *i* and *j* are both proper names, return true; else return false.
11. **Alias Feature (ALIAS):** Its possible values are true or false. If *i* is an alias of *j* or vice versa, return true; else return false. That is, this feature value is true if *i* and *j* are named entities (person, date, organization, etc.) that refer to the same entity. The alias module works differently depending

on the named entity type. For i and j that are dates (e.g., *01-08* and *Jan. 8*), by using string comparison, the day, month, and year values are extracted and compared. If they match, then j is an alias of i . For i and j that are “person,” such as *Mr. Simpson* and *Bert Simpson*, the last words of the noun phrases are compared to determine whether one is an alias of the other. For organization names, the alias function also checks for acronym match such as *IBM* and *International Business Machines Corp.* In this case, the longer string is chosen to be the one that is converted into the acronym form. The first step is to remove all postmodifiers such as *Corp.* and *Ltd.* Then, the acronym function considers each word in turn, and if the first letter is capitalized, it is used to form the acronym. Two variations of the acronyms are produced: one with a period after each letter, and one without.

12. **Appositive Feature (APPOSITIVE):** Its possible values are true or false. If j is in apposition to i , return true; else return false. For example, the markable *the chairman of Microsoft Corp.* is in apposition to *Bill Gates* in the sentence *Bill Gates, the chairman of Microsoft Corp., . . .*. Our system determines whether j is a possible appositive construct by first checking for the existence of verbs and proper punctuation. Like the above example, most appositives do not have any verb; and an appositive is separated by a comma from the most immediate antecedent, i , to which it refers. Further, at least one of i and j must be a proper name. The MUC-6 and MUC-7 coreference task definitions are slightly different. In MUC-6, j needs to be a definite noun phrase to be an appositive, while both indefinite and definite noun phrases are acceptable in MUC-7.

As an example, Table 1 shows the feature vector associated with the antecedent i , *Frank Newman*, and the anaphor j , *vice chairman*, in the following sentence:

- (1) Separately, Clinton transition officials said that *Frank Newman*, 50, *vice chairman* and chief financial officer of BankAmerica Corp., is expected to be nominated as assistant Treasury secretary for domestic finance.

Table 1

Feature vector of the markable pair (i = *Frank Newman*, j = *vice chairman*).

Feature	Value	Comments
DIST	0	i and j are in the same sentence
I_PRONOUN	–	i is not a pronoun
J_PRONOUN	–	j is not a pronoun
STR_MATCH	–	i and j do not match
DEF_NP	–	j is not a definite noun phrase
DEM_NP	–	j is not a demonstrative noun phrase
NUMBER	+	i and j are both singular
SEMCLASS	1	i and j are both persons (This feature has three values: false(0), true(1), unknown(2).)
GENDER	1	i and j are both males (This feature has three values: false(0), true(1), unknown(2).)
PROPER_NAME	–	Only i is a proper name
ALIAS	–	j is not an alias of i
APPOSITIVE	+	j is in apposition to i

Because of capitalization, markables in the headlines of MUC-6 and MUC-7 documents are always considered proper names even though some are not. Our system solves this inaccuracy by first preprocessing a headline to correct the capitalization before passing it into the pipeline of NLP modules. Only those markables in the headline that appear in the text body as proper names have their capitalization changed to match those found in the text body. All other headline markables are changed to lowercase.

2.3 Generating Training Examples

Consider a coreference chain A1 - A2 - A3 - A4 found in an annotated training document. Only pairs of noun phrases in the chain that are immediately adjacent (i.e., A1 - A2, A2 - A3, and A3 - A4) are used to generate the positive training examples. The first noun phrase in a pair is always considered the antecedent, while the second is the anaphor. On the other hand, negative training examples are extracted as follows. Between the two members of each antecedent-anaphor pair, there are other markables extracted by our language-processing modules that either are not found in any coreference chain or appear in other chains. Each of them is then paired with the anaphor to form a negative example. For example, if markables a, b, and B1 appear between A1 and A2, then the negative examples are a - A2, b - A2, and B1 - A2. Note that a and b do not appear in any coreference chain, while B1 appears in another coreference chain.

For an annotated noun phrase in a coreference chain in a training document, the same noun phrase must be identified as a markable by our pipeline of language-processing modules before this noun phrase can be used to form a feature vector for use as a training example. This is because the information necessary to derive a feature vector, such as semantic class and gender, is computed by the language-processing modules. If an annotated noun phrase is not identified as a markable, it will not contribute any training example. To see more clearly how training examples are generated, consider the following four sentences:

- Sentence 1
 1. (Eastern Air)_{a1} Proposes (Date For Talks on ((Pay)_{c1}-Cut)_{d1} Plan)_{b1}
 2. (Eastern Air)₁ Proposes (Date)₂ For (Talks)₃ on (Pay-Cut Plan)₄
- Sentence 2
 1. (Eastern Airlines)_{a2} executives notified (union)_{e1} leaders that the carrier wishes to discuss selective ((wage)_{c2} reductions)_{d2} on (Feb. 3)_{b2}.
 2. ((Eastern Airlines)₅ executives)₆ notified ((union)₇ leaders)₈ that (the carrier)₉ wishes to discuss (selective (wage)₁₀ reductions)₁₁ on (Feb. 3)₁₂.
- Sentence 3
 1. ((Union)_{e2} representatives who could be reached)_{f1} said (they)_{f2} hadn't decided whether (they)_{f3} would respond.
 2. ((Union)₁₃ representatives)₁₄ who could be reached said (they)₁₅ hadn't decided whether (they)₁₆ would respond.

- Sentence 4

1. By proposing (a meeting date)_{b3}, (Eastern)_{a3} moved one step closer toward reopening current high-cost contract agreements with ((its)_{a4} unions)_{e3}.
2. By proposing (a meeting date)₁₇, (Eastern)₁₈ moved (one step)₁₉ closer toward reopening (current high-cost contract agreements)₂₀ with ((its)₂₁ unions)₂₂.

Each sentence is shown twice with different noun phrase boundaries. Sentences labeled (1) are obtained directly from part of the training document. The letters in the subscripts uniquely identify the coreference chains, while the numbers identify the noun phrases. Noun phrases in sentences labeled (2) are extracted by our language-processing modules and are also uniquely identified by numeric subscripts.

Let's consider chain *e*, which is about the union. There are three noun phrases that corefer, and our system managed to extract the boundaries that correspond to all of them: (*union*)₇ matches with (*union*)_{e1}, (*union*)₁₃ with (*union*)_{e2}, and (*its unions*)₂₂ with (*its unions*)_{e3}. There are two positive training examples formed by ((*union*)₁₃, (*its unions*)₂₂) and ((*union*)₇, (*union*)₁₃). Noun phrases between (*union*)₇ and (*union*)₁₃ that do not corefer with (*union*)₁₃ are used to form the negative examples. The negative examples are ((*the carrier*)₉, (*union*)₁₃), ((*wage*)₁₀, (*union*)₁₃), ((*selective wage reductions*)₁₁, (*union*)₁₃), and ((*Feb. 3*)₁₂, (*union*)₁₃). Negative examples can also be found similarly between ((*union*)₁₃, (*its unions*)₂₂).

As another example, neither noun phrase in chain *d*, (*Pay-Cut*)_{d1} and (*wage reductions*)_{d2}, matches with any machine-extracted noun phrase boundaries. In this case, no positive or negative example is formed for noun phrases in chain *d*.

2.4 Building a Classifier

The next step is to use a machine learning algorithm to learn a classifier based on the feature vectors generated from the training documents. The learning algorithm used in our coreference engine is C5, which is an updated version of C4.5 (Quinlan 1993). C5 is a commonly used decision tree learning algorithm and thus it may be considered as a baseline method against which other learning algorithms can be compared.

2.5 Generating Coreference Chains for Test Documents

Before determining the coreference chains for a test document, all possible markables need to be extracted from the document. Every markable is a possible anaphor, and every markable before the anaphor in document order is a possible antecedent of the anaphor, except when the anaphor is nested. If the anaphor is a child or nested markable, then its possible antecedents must not be any markable with the same root markable as the current anaphor. However, the possible antecedents can be other root markables and their children that are before the anaphor in document order. For example, consider the two root markables, *Mr. Tom's daughter* and *His daughter's eyes*, appearing in that order in a test document. The possible antecedents of *His* cannot be *His daughter* or *His daughter's eyes*, but can be *Mr. Tom* or *Mr. Tom's daughter*.

The coreference resolution algorithm considers every markable *j* starting from the second markable in the document to be a potential candidate as an anaphor. For each *j*, the algorithm considers every markable *i* before *j* as a potential antecedent. For each pair *i* and *j*, a feature vector is generated and given to the decision tree classifier. A coreferencing antecedent is found if the classifier returns true. The algorithm starts from the immediately preceding markable and proceeds backward in the reverse order of

the markables in the document until there is no remaining markable to test or an antecedent is found.

As an example, consider the following text with markables already detected by the NLP modules:

- (2) **(Ms. Washington)**₇₃'s candidacy is being championed by (several powerful lawmakers)₇₄ including **((her)**₇₆ boss)₇₅, (Chairman John Dingell)₇₇ (D., (Mich.)₇₈) of (the House Energy and Commerce Committee)₇₉. **(She)**₈₀ currently is (a counsel)₈₁ to (the committee)₈₂. **(Ms. Washington)**₈₃ and (Mr. Dingell)₈₄ have been considered (allies)₈₅ of (the (securities)₈₇ exchanges)₈₆, while (banks)₈₈ and ((futures)₉₀ exchanges)₈₉ have often fought with (them)₉₁.

We will consider how the boldfaced chains are detected. Table 2 shows the pairs of markables tested for coreference to form the chain for *Ms. Washington-her-She-Ms. Washington*. When the system considers the anaphor, *(her)*₇₆, all preceding phrases, except *(her boss)*₇₅, are tested to see whether they corefer with it. *(her boss)*₇₅ is not tested because *(her)*₇₆ is its nested noun phrase. Finally, the decision tree determines that the noun phrase *(Ms. Washington)*₇₃ corefers with *(her)*₇₆. In Table 2, we only show the system considering the three anaphors *(her)*₇₆, *(She)*₈₀, and *(Ms. Washington)*₈₃, in that order.

Table 2

Pairs of markables that are tested in forming the coreference chain *Ms. Washington-her-She-Ms. Washington*. The feature vector format: DIST, SEMCLASS, NUMBER, GENDER, PROPER_NAME, ALIAS, J_PRONOUN, DEF_NP, DEM_NP, STR_MATCH, APPOSITIVE, I_PRONOUN.

Antecedent	Anaphor	Feature Vector	Corefers?
(several powerful lawmakers) ₇₄	(her) ₇₆	0,1,-2,-,-,+,-,-,-,-	No
(Ms. Washington) ₇₃	(her) ₇₆	0,1,+1,-,-,+,-,-,-,-	Yes
(the House Energy and Commerce Committee) ₇₉	(She) ₈₀	1,0,+0,-,-,+,-,-,-,-	No
(Mich.) ₇₈	(She) ₈₀	2,0,+0,-,-,+,-,-,-,-	No
(Chairman John Dingell) ₇₇	(She) ₈₀	3,1,+0,-,-,+,-,-,-,-	No
(her) ₇₆	(She) ₈₀	3,1,+1,-,-,+,-,-,-,-	Yes
(the committee) ₈₂	(Ms. Washington) ₈₃	1,0,+0,-,-,-,-,-,-,-	No
(a counsel) ₈₁	(Ms. Washington) ₈₃	1,1,+2,-,-,-,-,-,-,-	No
(She) ₈₀	(Ms. Washington) ₈₃	1,1,+1,-,-,-,-,-,-,-	No
(the House Energy and Commerce Committee) ₇₉	(Ms. Washington) ₈₃	2,0,+0,+,-,-,-,-,-,-	No
(Mich.) ₇₈	(Ms. Washington) ₈₃	3,0,+0,+,-,-,-,-,-,-	No
(Chairman John Dingell) ₇₇	(Ms. Washington) ₈₃	4,1,+0,+,-,-,-,-,-,-	No
(her) ₇₆	(Ms. Washington) ₈₃	4,1,+1,-,-,-,-,-,-,-	No
(her boss) ₇₅	(Ms. Washington) ₈₃	4,1,-0,-,-,-,-,-,-,-	No
(several powerful lawmakers) ₇₄	(Ms. Washington) ₈₃	4,1,-2,-,-,-,-,-,-,-	No
(Ms. Washington) ₇₃	(Ms. Washington) ₈₃	4,1,+1,+,-,-,-,-,-,-	Yes

We use the same method to generate coreference chains for both MUC-6 and MUC-7, except for the following. For MUC-7, because of slight changes in the coreference task definition, we include a filtering module to remove certain coreference chains. The task definition states that a coreference chain must contain at least one element that is a head noun or a name; that is, a chain containing only prenominal modifiers is removed by the filtering module.

3. Evaluation

In order to evaluate the performance of our learning approach to coreference resolution on common data sets, we utilized the annotated corpora and scoring programs from MUC-6 and MUC-7, which assembled a set of newswire documents annotated with coreference chains. Although we did not participate in either MUC-6 or MUC-7, we were able to obtain the training and test corpora for both years from the MUC organizers for research purposes.¹ To our knowledge, these are the only publicly available annotated corpora for coreference resolution.

For MUC-6, 30 dry-run documents annotated with coreference information were used as the training documents for our coreference engine. There are also 30 annotated training documents from MUC-7. The total size of the 30 training documents is close to 12,400 words for MUC-6 and 19,000 words for MUC-7. There are altogether 20,910 (48,872) training examples used for MUC-6 (MUC-7), of which only 6.5% (4.4%) are positive examples in MUC-6 (MUC-7).²

After training a separate classifier for each year, we tested the performance of each classifier on its corresponding test corpus. For MUC-6, the C5 pruning confidence is set at 20% and the minimum number of instances per leaf node is set at 5. For MUC-7, the pruning confidence is 60% and the minimum number of instances is 2. The parameters are determined by performing 10-fold cross-validation on the whole training set for each MUC year. The possible pruning confidence values that we tried are 10%, 20%, 40%, 60%, 80%, and 100%, and for minimum instances, we tried 2, 5, 10, 15, and 20. Thus, a total of 30 (6×5) cross-validation runs were executed.

One advantage of using a decision tree learning algorithm is that the resulting decision tree classifier can be interpreted by humans. The decision tree generated for MUC-6, shown in Figure 2, seems to encapsulate a reasonable rule of thumb that matches our intuitive linguistic notion of when two noun phrases can corefer. It is also interesting to note that only 8 out of the 12 available features in the training examples are actually used in the final decision tree built.

MUC-6 has a standard set of 30 test documents, which is used by all systems that participated in the evaluation. Similarly, MUC-7 has a test corpus of 20 documents. We compared our system's MUC-6 and MUC-7 performance with that of the systems that took part in MUC-6 and MUC-7, respectively. When the coreference engine is given new test documents, its output is in the form of SGML files with the coreference chains properly annotated according to the guidelines.³ We then used the scoring programs

1 See http://www.itl.nist.gov/iad/894.02/related_projects/muc/index.html for details on obtaining the corpora.

2 Our system runs on a Pentium III 550MHz PC. It took less than 5 minutes to generate the training examples from the training documents for MUC-6, and about 7 minutes for MUC-7. The training time for the C5 algorithm to generate a decision tree from all the training examples was less than 3 seconds for both MUC years.

3 The time taken to generate the coreference chains for the 30 MUC-6 test documents of close to 13,400 words was less than 3 minutes, while it took less than 2 minutes for the 20 MUC-7 test documents of about 10,000 words.

```

STR_MATCH = +: +
STR_MATCH = -:
:...J_PRONOUN = -:
    :...APPOSITIVE = +: +
    :    APPOSITIVE = -:
    :    :...ALIAS = +: +
    :    ALIAS = -: -
J_PRONOUN = +:
:...GENDER = 0: -
    GENDER = 2: -
    GENDER = 1:
    :...I_PRONOUN = +: +
    I_PRONOUN = -:
    :...DIST > 0: -
    DIST <= 0:
    :...NUMBER = +: +
    NUMBER = -: -

```

Figure 2

The decision tree classifier learned for MUC-6.

for the respective years to generate the recall and precision scores for our coreference engine.

Our coreference engine achieves a recall of 58.6% and a precision of 67.3%, yielding a balanced F-measure of 62.6% for MUC-6. For MUC-7, the recall is 56.1%, the precision is 65.5%, and the balanced F-measure is 60.4%.⁴ We plotted the scores of our coreference engine (square-shaped) against the official test scores of the other systems (cross-shaped) in Figure 3 and Figure 4.

We also plotted the learning curves of our coreference engine in Figure 5 and Figure 6, showing its accuracy averaged over three random trials when trained on 1, 2, 3, 4, 5, 10, 15, 20, 25, and 30 training documents. The learning curves indicate that our coreference engine achieves its peak performance with about 25 training documents, or about 11,000 to 17,000 words of training documents. This number of training documents would generate tens of thousands of training examples, sufficient for the decision tree learning algorithm to learn a good classifier. At higher numbers of training documents, our system seems to start overfitting the training data. For example, on MUC-7 data, training on the full set of 30 training documents results in a more complex decision tree.

Our system's scores are in the upper region of the MUC-6 and MUC-7 systems. We performed a simple one-tailed, paired sample t-test at significance level $p = 0.05$ to determine whether the difference between our system's F-measure score and each of the other systems' F-measure score on the test documents is statistically significant.⁵ We found that at the 95% significance level ($p = 0.05$), our system performed better than three MUC-6 systems, and as well as the rest of the MUC-6 systems. Using the

⁴ Note that MUC-6 did not use balanced F-measure as the official evaluation measure, but MUC-7 did.

⁵ Though the McNemar test is shown to have low Type I error compared with the paired t-test (Dietterich 1998), we did not carry out this test in the context of coreference. This is because an example instance defines a coreference link between two noun phrases, and since this link is transitive in nature, it is unclear how the number of links misclassified by System A but not by System B and vice versa can be obtained to execute the McNemar test.




Figure 3
Coreference scores of MUC-6 systems and our system.




Figure 4
Coreference scores of MUC-7 systems and our system.

same significance level, our system performed better than four MUC-7 systems, and as well as the rest of the MUC-7 systems. Our result is encouraging since it indicates that a learning approach using relatively shallow features can achieve scores comparable to those of systems built using nonlearning approaches.




Figure 5
Learning curve of coreference resolution accuracy for MUC-6.

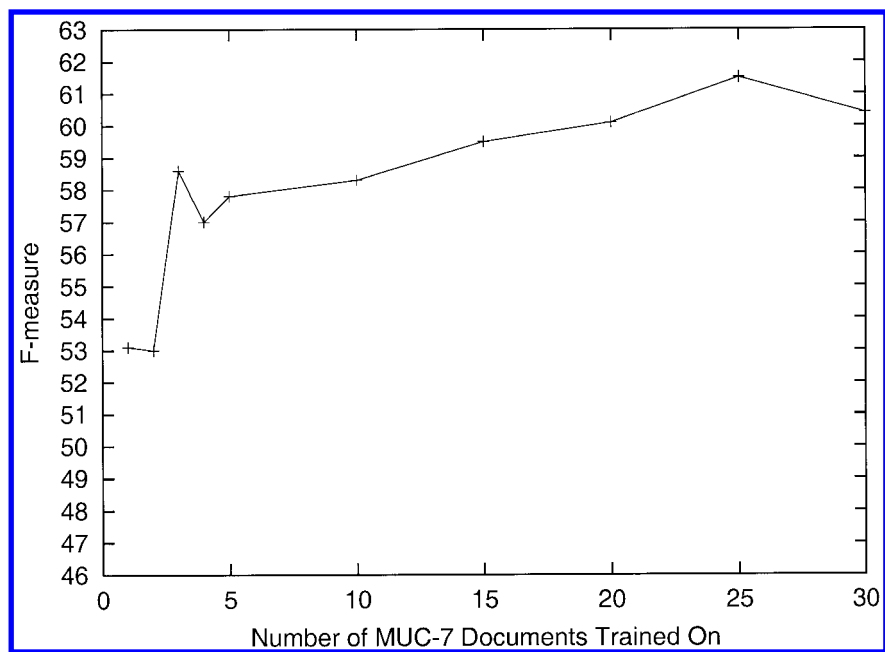


Figure 6
Learning curve of coreference resolution accuracy for MUC-7.

It should be noted that the accuracy of our coreference resolution engine depends to a large extent on the performance of the NLP modules that are executed before the coreference engine. Our current learning-based, HMM named entity recognition module is trained on 318 documents (a disjoint set from both the MUC-6 and MUC-7

test documents) tagged with named entities, and its score on the MUC-6 named entity task for the 30 formal test documents is only 88.9%, which is not considered very high by MUC-6 standards. For example, our named entity recognizer could not identify the two named entities *USAir* and *Piedmont* in the expression *USAir and Piedmont* but instead treat them as one single named entity. Our part-of-speech tagger achieves 96% accuracy, while the accuracy of noun phrase identification is above 90%.

4. The Contribution of the Features

One factor that affects the performance of a machine learning approach is the set of features used. It is interesting to find out how useful each of our 12 features is in the MUC-6 and MUC-7 coreference tasks. One way to do this is to train and test using just one feature at a time. Table 3 and Table 4 show the results of the experiment. For both MUC-6 and MUC-7, the 3 features that give nonzero recall and precision are ALIAS, STR_MATCH, and APPOSITIVE. The 12 features can be divided into unary and binary

Table 3

MUC-6 results of complete and baseline systems to study the contribution of the features.

System ID	Recall	Prec	F	Remarks
Complete systems				
DSO	58.6	67.3	62.6	Our system
DSO_TRG	52.6	67.6	59.2	Our system using RESOLVE's method of generating positive and negative examples
RESOLVE	44.2	50.7	47.2	The RESOLVE coreference system at the University of Massachusetts
Baseline systems using just one feature				
DIST	0.0	0.0	0.0	Only "distance" feature is used
SEMCLASS	0.0	0.0	0.0	Only "semantic class agreement"
NUMBER	0.0	0.0	0.0	Only "number agreement"
GENDER	0.0	0.0	0.0	Only "gender agreement"
PROPER_NAME	0.0	0.0	0.0	Only "both proper names"
ALIAS	24.5	88.7	38.4	Only "alias"
J_PRONOUN	0.0	0.0	0.0	Only "j-pronoun"
DEF_NP	0.0	0.0	0.0	Only "definite noun phrase"
DEM_NP	0.0	0.0	0.0	Only "demonstrative noun phrase"
STR_MATCH	45.7	65.6	53.9	Only "string match"
APPOSITIVE	3.9	57.7	7.3	Only "appositive"
I_PRONOUN	0.0	0.0	0.0	Only "i-pronoun"
Other baseline systems				
ALIAS_STR	51.5	66.4	58.0	Only the "alias" and "string match" features are used
ALIAS_STR_APPOS	55.2	66.4	60.3	Only the "alias," "string match," and "appositive" features are used
ONE_CHAIN	89.9	31.8	47.0	All markables form one chain
ONE_WRD	55.4	36.6	44.1	Markables corefer if there is at least one common word
HD_WRD	56.4	50.4	53.2	Markables corefer if their head words are the same

Table 4

MUC-7 results of complete and baseline systems to study the contribution of the features.

System ID	Recall	Prec	F	Remarks
Complete systems				
DSO	56.1	65.5	60.4	Our system
DSO_TRG	53.3	69.7	60.4	Our system using RESOLVE's method of generating positive and negative examples
Baseline systems using just one feature				
DIST	0.0	0.0	0.0	Only "distance" feature is used
SEMCLASS	0.0	0.0	0.0	Only "semantic class agreement"
NUMBER	0.0	0.0	0.0	Only "number agreement"
GENDER	0.0	0.0	0.0	Only "gender agreement"
PROPER_NAME	0.0	0.0	0.0	Only "both proper names"
ALIAS	25.6	81.1	38.9	Only "alias"
J_PRONOUN	0.0	0.0	0.0	Only "j-pronoun"
DEF_NP	0.0	0.0	0.0	Only "definite noun phrase"
DEM_NP	0.0	0.0	0.0	Only "demonstrative noun phrase"
STR_MATCH	43.8	71.4	54.3	Only "string match"
APPOSITIVE	2.4	60.0	4.6	Only "appositive"
LPRONOUN	0.0	0.0	0.0	Only "i-pronoun"
Other baseline systems				
ALIAS_STR	49.4	70.4	58.1	Only the "alias" and "string match" features are used
ALIAS_STR_APPOS	51.6	69.9	59.4	Only the "alias," "string match," and "appositive" features are used
ONE_CHAIN	87.5	30.5	45.2	All markables form one chain
ONE_WRD	55.9	38.7	45.7	Markables corefer if there is at least one common word
HD_WRD	55.2	55.6	55.4	Markables corefer if their head words are the same

features. The unary features are LPRONOUN, J_PRONOUN, DEF_NP, and DEM_NP, while the rest are binary in nature. All the unary features score an F-measure of 0. The binary features with 0 F-measure are DIST, PROPER_NAME, GENDER, SEMCLASS, and NUMBER.

The ALIAS, APPOSITIVE, and STR_MATCH features give nonzero F-measure. All these features give rather high precision scores ($> 80\%$ for ALIAS, $> 65\%$ for STR_MATCH, and $> 57\%$ for APPOSITIVE). Since these features are highly informative, we were curious to see how much they contribute to our MUC-6 and MUC-7 results of 62.6% and 60.4%, respectively. Systems ALIAS_STR and ALIAS_STR_APPOS in Table 3 and Table 4 show the results of the experiment. In terms of absolute F-measure, the difference between using these three features and using all features is 2.3% for MUC-6 and 1% for MUC-7; in other words, the other nine features contribute just 2.3% and 1% more for each of the MUC years. These nine features will be the first ones to be considered for pruning away by the C5 algorithm. For example, four features, namely, SEMCLASS, PROPER_NAME, DEF_NP, and DEM_NP, are not used in the MUC-6 tree shown in Figure 2. Figure 7 shows the distribution of the test cases over the five positive leaf nodes of the MUC-6 tree. For example, about 66.3% of all

```

STR_MATCH = +: + 944 (66.3%)

STR_MATCH = -:
...J_PRONOUN = -:
...APPOSITIVE = +: + 111 (7.8%)

STR_MATCH = -:
...J_PRONOUN = -:
...APPOSITIVE = -:
...ALIAS = +: + 163 (11.5%)

STR_MATCH = -:
...J_PRONOUN = +:
...GENDER = 1:
...I_PRONOUN = +: + 77 (5.4%)

STR_MATCH = -:
...J_PRONOUN = +:
...GENDER = 1:
...I_PRONOUN = -:
...DIST <= 0:
...NUMBER = +: + 128 (9.0%)

```

Figure 7

Distribution of test examples from the 30 MUC-6 test documents for positive leaf nodes of the MUC-6 tree.

the test examples that are classified positive go to the “If STR_MATCH” branch of the tree.

Other baseline systems that are used are ONE_CHAIN, ONE_WRD, and HD_WRD (Cardie and Wagstaff 1999). For ONE_CHAIN, all markables formed one chain. In ONE_WRD, markables corefer if there is at least one common word. In HD_WRD, markables corefer if their head words are the same. The purpose of ONE_CHAIN is to determine the maximum recall our system is capable of. The recall level here indirectly measures how effective the noun phrase identification module is. Both ONE_WRD and HD_WRD are less stringent variations of STR_MATCH. The performance of ONE_WRD is the worst. HD_WRD offers better recall compared to STR_MATCH, but poorer precision. However, its F-measure is comparable to that of STR_MATCH.

The score of the coreference system at the University of Massachusetts (RESOLVE), which uses C4.5 for coreference resolution, is shown in Table 3. RESOLVE is shown because among the MUC-6 systems, it is the only machine learning-based system that we can directly compare to. The other MUC-6 systems were not based on a learning approach. Also, none of the systems in MUC-7 adopted a learning approach to coreference resolution (Chinchor 1998).

RESOLVE’s score is not high compared to scores attained by the rest of the MUC-6 systems. In particular, the system’s recall is relatively low. Our system’s score is higher than that of RESOLVE, and the difference is statistically significant. The RESOLVE system is described in three papers: McCarthy and Lehnert (1995), Fisher et al. (1995), and McCarthy (1996). As explained in McCarthy (1996), the reason for this low recall is that RESOLVE takes only the “relevant entities” and “relevant references” as input, where the relevant entities and relevant references are restricted to “person”

and “organization.” In addition, because of limitations of the noun phrase detection module, nested phrases are not extracted and therefore do not take part in coreference. Nested phrases can include prenominal modifiers, possessive pronouns, and so forth. Therefore, the number of candidate markables to be used for coreference is small.

On the other hand, the markables extracted by our system include nested noun phrases, MUC-style named entity types (money, percent, date, etc.), and other types not defined by MUC. These markables will take part in coreference. About 3,600 top-level markables are extracted from the 30 MUC-6 test documents by our system. As detected by our NLP modules, only about 35% of these 3,600 phrases are “person” and “organization” entities and references. Concentrating on just these types has thus affected the overall recall of the RESOLVE system.

RESOLVE’s way of generating training examples also differs from our system’s: instances are created for all possible pairings of “relevant entities” and “relevant references,” instead of our system’s method of stopping at the first coreferential noun phrase when traversing back from the anaphor under consideration. We implemented RESOLVE’s way of generating training examples, and the results (DSO-TRG) are reported in Table 3 and Table 4. For MUC-7, there is no drop in F-measure; for MUC-6, the F-measure dropped slightly.

RESOLVE makes use of 39 features, considerably more than our system’s 12 features. RESOLVE’s feature set includes the two highly informative features, ALIAS and STR.MATCH. RESOLVE does not use the APPOSITIVE feature.

5. Error Analysis

In order to determine the major classes of errors made by our system, we randomly chose five test documents from MUC-6 and determined the coreference links that were either missing (false negatives) or spurious (false positives) in these sample documents. Missing links result in recall errors; spurious links result in precision errors.

Breakdowns of the number of spurious and missing links are shown in Table 5 and Table 6, respectively. The following two subsections describe the errors in more detail.

5.1 Errors Causing Spurious Links

This section describes the five major types of errors summarized in Table 5 in more detail.

5.1.1 Prenominal Modifier String Match. This class of errors occurs when some strings of the prenominal modifiers of two markables match by surface string comparison and thus, by the C5 decision tree in Figure 2, the markables are treated as coreferring. How-

Table 5
The types and frequencies of errors that affect precision.

Types of Errors Causing Spurious Links	Frequency	%
Prenominal modifier string match	16	42.1%
Strings match but noun phrases refer to different entities	11	28.9%
Errors in noun phrase identification	4	10.5%
Errors in apposition determination	5	13.2%
Errors in alias determination	2	5.3%

Table 6

The types and frequencies of errors that affect recall.

Types of Errors Causing Missing Links	Frequency	%
Inadequacy of current surface features	38	63.3%
Errors in noun phrase identification	7	11.7%
Errors in semantic class determination	7	11.7%
Errors in part-of-speech assignment	5	8.3%
Errors in apposition determination	2	3.3%
Errors in tokenization	1	1.7%

ever, the entire markable actually does not corefer. The nested noun phrase extraction module is responsible for obtaining the possible prenominal modifiers from a noun phrase.

In (3), the noun phrase extraction module mistakenly extracted *(vice)₁* and *(vice)₂*, which are not prenominal modifiers. Because of string match, *(vice)₁* and *(vice)₂* incorrectly corefer. In (4), *(undersecretary)₂* was correctly extracted as a prenominal modifier, but incorrectly corefers with *(undersecretary)₁* by string match.

- (3) David J. Bronczek, **(vice)₁** president and general manager of Federal Express Canada Ltd., was named senior **(vice)₂** president, Europe, Africa and Mediterranean, at this air-express concern.
- (4) Tarnoff, a former Carter administration official and president of the Council on Foreign Relations, is expected to be named **(undersecretary)₁** for political affairs. . . . Former Sen. Tim Wirth is expected to get a newly created **(undersecretary)₂** post for global affairs, which would include refugees, drugs and environmental issues.

5.1.2 Strings Match but Noun Phrases Refer to Different Entities. This error occurs when the surface strings of two markables match and thus, by the C5 decision tree in Figure 2, they are treated as coreferring. However, they actually refer to different entities and should not corefer. In (5), *(the committee)₁* actually refers to the entity *the House Energy and Commerce Committee*, and *(the committee)₂* refers to *the Senate Finance Committee*; therefore, they should not corefer. In (6), the two instances of *chief executive officer* refer to two different persons, namely, *Allan Laufgraben* and *Milton Petrie*, and, again, should not corefer.

- (5) Ms. Washington's candidacy is being championed by several powerful lawmakers including her boss, Chairman John Dingell (D., Mich.) of the House Energy and Commerce Committee. She currently is a counsel to **(the committee)₁**. . . . Mr. Bentsen, who headed the Senate Finance Committee for the past six years, also is expected to nominate Samuel Sessions, **(the committee)₂**'s chief tax counsel, to one of the top tax jobs at Treasury.
- (6) Directors also approved the election of Allan Laufgraben, 54 years old, as president and **(chief executive officer)₁** and Peter A. Left, 43, as chief operating officer. Milton Petrie, 90-year-old chairman, president and **(chief executive officer)₂** since the company was founded in 1932, will continue as chairman.

5.1.3 Errors in Noun Phrase Identification. This class of errors is caused by mistakes made by the noun phrase identification module. In (7), *May* and *June* are incorrectly grouped together by the noun phrase identification module as one noun phrase, that is, *May, June*. This markable then incorrectly causes the APPOSITIVE feature to be true, which results in classifying the pair as coreferential. In fact, *(the first week of July)*₂ should not be in apposition to *(May, June)*₁. However, we classified this error as a noun phrase identification error because it is the first module that causes the error. In (8), the noun phrase module extracted *Metaphor Inc.* instead of *Metaphor Inc. unit*. This causes *(it)*₂ to refer to *Metaphor Inc.* instead of *Metaphor Inc. unit*.

- (7) The women's apparel specialty retailer said sales at stores open more than one year, a key barometer of a retail concern's strength, declined 2.5% in **(May, June)**₁ and **(the first week of July)**₂.
- (8) ... International Business Machines Corp.'s **(Metaphor Inc.)**₁ unit said **(it)**₂ will shed 80 employees ...

5.1.4 Errors in Apposition Determination. This class of errors occurs when the anaphor is incorrectly treated as being in apposition to the antecedent and therefore causes the noun phrases to corefer. The precision scores obtained when using the APPOSITIVE feature alone are shown in Table 3 and Table 4, which suggest that the module can be improved further. Examples where apposition determination is incorrect are shown in (9) and (10).

- (9) Clinton officials are said to be deciding between recently retired Rep. Matthew McHugh (D., **(N.Y.)**₁) and **(environmental activist)**₂ and transition official Gus Speth for the director of the Agency for International Development.
- (10) Metaphor, a software subsidiary that IBM purchased in 1991, also named **(Chris Grejtak)**₁, **(43 years old)**₂, currently a senior vice president, president and chief executive officer.

5.1.5 Errors in Alias Determination. This class of errors occurs when the anaphor is incorrectly treated as an alias of the antecedent, thus causing the noun phrase pair to corefer. In (11), the two phrases *(House)*₁ and *(the House Energy and Commerce Committee)*₂ corefer because the ALIAS feature is incorrectly determined to be true.

- (11) Consuela Washington, a longtime **(House)**₁ staffer and an expert in securities laws, is a leading candidate to be chairwoman of the Securities and Exchange Commission in the Clinton administration. ... Ms. Washington's candidacy is being championed by several powerful lawmakers including her boss, Chairman John Dingell (D., Mich.) of **(the House Energy and Commerce Committee)**₂.

5.2 Errors Causing Missing Links

This subsection describes the six major classes of errors summarized in Table 6 in more detail.

5.2.1 Inadequacy of Current Surface Features. This class of errors is due to the inadequacy of the current surface features because they do not have information about

other words (such as the connecting conjunctions, prepositions, or verbs) and other knowledge sources that may provide important clues for coreference. As a result, the set of shallow features we used is unable to correctly classify the noun phrases in the examples below as coreferring.

Example (12) illustrates why resolving *(them)*₂ is difficult. *(allies)*₁, *securities exchanges*, *banks*, and *futures exchanges* are all possible antecedents of *(them)*₂, and the feature set must include more information to be able to pick the correct one. The conjunction *and* in (13) and *was named* in (16) are important cues to determine coreference. In addition, it may also be possible to capture noun phrases in predicate constructions like (17), where *(Mr. Gleason)*₁ is the subject and *(president)*₂ is the object.

- (12) Ms. Washington and Mr. Dingell have been considered **(allies)**₁ of the securities exchanges, while banks and futures exchanges have often fought with **(them)**₂.
- (13) Separately, Clinton transition officials said that Frank Newman, 50, **(vice chairman)**₁ and **(chief financial officer)**₂ of BankAmerica Corp., is expected to be nominated as assistant Treasury secretary for domestic finance.
- (14) Separately, **(Clinton transition officials)**₁ said that Frank Newman, 50, vice chairman and chief financial officer of BankAmerica Corp., is expected to be nominated as assistant Treasury secretary for domestic finance. . . . As early as today, **(the Clinton camp)**₂ is expected to name five undersecretaries of state and several assistant secretaries.
- (15) **(Metro-Goldwyn-Mayer Inc.)**₁ said it named Larry Gleason president of world-wide theatrical distribution of **(the movie studio)**₂'s distribution unit.
- (16) . . . **(general manager)**₁ of Federal Express Canada Ltd., was named **(senior vice president)**₂, Europe, Africa and Mediterranean . . .
- (17) **(Mr. Gleason)**₁, 55 years old, was **(president)**₂ of theatrical exhibition for Paramount Communications Inc., in charge of the company's 1,300 movie screens in 12 countries.

5.2.2 Errors in Noun Phrase Identification. This class of errors was described in Section 5.1.3. The noun phrase identification module may extract noun phrases that do not match the phrases in the coreference chain, therefore causing missing links and recall error.

5.2.3 Errors in Semantic Class Determination. These errors are caused by the wrong assignment of semantic classes to words. For example, *(Metaphor)*₁ should be assigned "organization" but it is assigned "unknown" in (18), and *(second-quarter)*₂ should be assigned "date" instead of "unknown" in (19). However, correcting these classes will still not cause the noun phrases in the examples to corefer. This is because the values of the SEMCLASS feature in the training examples are extremely noisy, a situation caused largely by our semantic class determination module. In many of the negative training examples, although the noun phrases are assigned the same semantic classes, these assignments do not seem to be correct. Some examples are *(four-year, NBC)*, *(The union, Ford Motor Co.)*, and *(base-wage, job-security)*. A better algorithm for assigning semantic classes and a more refined semantic class hierarchy are needed.

- (18) **(Metaphor)**₁, a software subsidiary that IBM purchased in 1991, also named Chris Grejtak, . . . Mr. Grejtak said in an interview that the staff reductions will affect most areas of **(the company)**₂ related to its early proprietary software products.
- (19) Business brief—Petrie Stores Corp.: Losses for **(Fiscal 2nd Period)**₁, half seen likely by retailer. . . . Petrie Stores Corp., Secaucus, N.J., said an uncertain economy and faltering sales probably will result in a **(second-quarter)**₂ loss and perhaps a deficit for the first six months of fiscal 1994.

5.2.4 Errors in Part-of-Speech Assignment. This class of errors is caused by the wrong assignment of part-of-speech tags to words. In (20), *(there)*₂ is not extracted because the part-of-speech tag assigned is “RB,” which is an adverb and not a possible noun phrase.

- (20) Jon W. Slingerup, who is 43 and has been director of customer service in **(Canada)**₁, succeeds Mr. Bronczek as vice president and general manager **(there)**₂.

5.2.5 Errors in Apposition Determination. This class of errors was described in Section 5.1.4.

5.2.6 Errors in Tokenization. This class of errors is due to the incorrect tokenization of words. In (21), *(1-to-2)*₁ and *(15-to-1)*₂ are not found because the tokenizer breaks *1-to-2* into *1*, *-*, *to-2*. *15-to-1* is broken up similarly.

- (21) Separately, MGM said it completed a previously announced financial restructuring designed to clean up its balance sheet—removing \$900 million in bank debt from MGM’s books and reducing its debt-to-equity ratio to **(1-to-2)**₁ from **(15-to-1)**₂—with a view toward a future sale of the company.

5.3 Comparing Errors Made by RESOLVE

McCarthy (1996) has also performed an analysis of errors while conducting an evaluation on the MUC-5 English Joint Venture (EJV) corpus. A large number of the spurious links are caused by what he terms “feature ambiguity,” which means that feature values are not computed perfectly. As seen in Table 5, our string match feature accounts for most of the spurious links. Also, seven of the spurious links are caused by alias and apposition determination. As with RESOLVE, “feature ambiguity” is the main source of precision errors.

For RESOLVE, a large number of the missing links are caused by “incomplete semantic knowledge” (32%) and “unused features” (40.5%). For our system, the errors due to the inadequacy of surface features and semantic class determination problems account for about 75% of the missing links. “Unused features” means that some of the features, or combinations of features, that are needed to classify pairs of phrases as coreferential are not present in the decision trees (McCarthy 1996). Similarly, the inadequacy of our system’s surface features means that the current feature set may not be enough and more information sources should be added.

Because a detailed error analysis of RESOLVE would require not only its MUC-6 response file, but also the output of its various components, we cannot perform the same error analysis that we did for our system on RESOLVE.

6. Related Work

There is a long tradition of work on coreference resolution within computational linguistics, but most of it was not subject to empirical evaluation until recently. Among the papers that have reported quantitative evaluation results, most are not based on learning from an annotated corpus (Baldwin 1997; Kameyama 1997; Lappin and Leass 1994; Mitkov 1997).

To our knowledge, the research efforts of Aone and Bennett (1995), Ge, Hale, and Charniak (1998), Kehler (1997), McCarthy and Lehnert (1995), Fisher et al. (1995), and McCarthy (1996) are the only ones that are based on learning from an annotated corpus. Ge, Hale, and Charniak (1998) used a statistical model for resolving pronouns, whereas we used a decision tree learning algorithm and resolved general noun phrases, not just pronouns. Similarly, Kehler (1997) used maximum entropy modeling to assign a probability distribution to alternative sets of coreference relationships among noun phrase entity templates, whereas we used decision tree learning.

The work of Aone and Bennett (1995), McCarthy and Lehnert (1995), Fisher et al. (1995), and McCarthy (1996) employed decision tree learning. The RESOLVE system is presented in McCarthy and Lehnert (1995), Fisher et al. (1995), and McCarthy (1996). McCarthy and Lehnert (1995) describe how RESOLVE was tested on the MUC-5 English Joint Ventures (EJV) corpus. It used a total of 8 features, 3 of which were specific to the EJV domain. For example, the feature *JV-CHILD-*i** determined whether *i* referred to a joint venture formed as the result of a tie-up. McCarthy (1996) describes how the original RESOLVE for MUC-5 EJV was improved to include more features, 8 of which were domain specific, and 30 of which were domain independent. Fisher et al. (1995) adapted RESOLVE to work in MUC-6. The features used were slightly changed for this domain. Of the original 30 domain-independent features, 27 were used. The 8 domain-specific features were completely changed for the MUC-6 task. For example, *JV-CHILD-*i** was changed to *CHILD-*i** to decide whether *i* is a “unit” or a “subsidiary” of a certain parent company. In contrast to RESOLVE, our system makes use of a smaller set of 12 features and, as in Aone and Bennett’s (1995) system, the features used are generic and applicable across domains. This makes our coreference engine a domain-independent module.

Although Aone and Bennett’s (1995) system also made use of decision tree learning for coreference resolution, they dealt with Japanese texts, and their evaluation focused only on noun phrases denoting organizations, whereas our evaluation, which dealt with English texts, encompassed noun phrases of all types, not just those denoting organizations. In addition, Aone and Bennett evaluated their system on noun phrases that had been correctly identified, whereas we evaluated our coreference resolution engine as part of a total system that first has to identify all the candidate noun phrases and has to deal with the inevitable noisy data when mistakes occur in noun phrase identification and semantic class determination.

The contribution of our work lies in showing that a learning approach, when evaluated on common coreference data sets, is able to achieve accuracy competitive with that of state-of-the-art systems using nonlearning approaches. It is also the first machine learning-based system to offer performance comparable to that of nonlearning approaches.

Finally, the work of Cardie and Wagstaff (1999) also falls under the machine learning approach. However, they used unsupervised learning and their method did not require any annotated training data. Their clustering method achieved a balanced F-measure of only 53.6% on MUC-6 test data. This is to be expected: supervised learning in general outperforms unsupervised learning since a supervised learning algorithm

has access to a richer set of annotated data to learn from. Since our supervised learning approach requires only a modest number of annotated training documents to achieve good performance (as can be seen from the learning curves), we argue that the better accuracy achieved more than justifies the annotation effort incurred.

7. Conclusion

In this paper, we presented a learning approach to coreference resolution of noun phrases in unrestricted text. The approach learns from a small, annotated corpus and the task includes resolving not just pronouns but general noun phrases. We evaluated our approach on common data sets, namely, the MUC-6 and MUC-7 coreference corpora. We obtained encouraging results, indicating that on the general noun phrase coreference task, the learning approach achieves accuracy comparable to that of non-learning approaches.

Acknowledgments

This paper is an expanded version of a preliminary paper that appeared in the *Proceedings of the 1999 Joint SIGDAT Conference on Empirical Methods in Natural Language Processing and Very Large Corpora*. We would like to thank the MUC organizers who made available to us the MUC-6 and MUC-7 data sets, without which this work would have been impossible. We also thank Beth Sundheim for helpful comments on an earlier version of this paper, and Hai Leong Chieu for his implementation of the HMM-based named entity recognition module.

References

- Aone, Chinatsu and Scott William Bennett. 1995. Evaluating automated and manual acquisition of anaphora resolution strategies. In *Proceedings of the 33rd Annual Meeting of the Association for Computational Linguistics*, 122–129.
- Baldwin, Breck. 1997. CogNIAC: High precision coreference with limited knowledge and linguistic resources. In *Proceedings of the ACL Workshop on Operational Factors in Practical, Robust Anaphora Resolution for Unrestricted Texts*, pages 38–45.
- Bikel, Daniel M., Richard Schwartz, and Ralph M. Weischedel. 1999. An algorithm that learns what's in a name. *Machine Learning*, 34(1–3):211–231, February.
- Cardie, Claire and Kiri Wagstaff. 1999. Noun phrase coreference as clustering. In *Proceedings of the 1999 Joint SIGDAT Conference on Empirical Methods in Natural Language Processing and Very Large Corpora*, pages 82–89.
- Chinchor, Nancy A. 1998. Overview of MUC-7/MET-2. In *Proceedings of the Seventh Message Understanding Conference (MUC-7)*. http://www.itl.nist.gov/iad/894.02/related_projects/muc/proceedings/muc_7_toc.html.
- Church, Kenneth. 1988. A stochastic parts program and noun phrase parser for unrestricted text. In *Proceedings of the Second Conference on Applied Natural Language Processing*, pages 136–143.
- Dietterich, Thomas G. 1998. Approximate statistical tests for comparing supervised classification learning algorithms. *Neural Computation*, 10(7):1895–1924, October.
- Fisher, David, Stephen Soderland, Joseph McCarthy, Fangfang Feng, and Wendy Lehnert. 1995. Description of the UMass system as used for MUC-6. In *Proceedings of the Sixth Message Understanding Conference (MUC-6)*, pages 127–140.
- Ge, Niyu, John Hale, and Eugene Charniak. 1998. A statistical approach to anaphora resolution. In *Proceedings of the Sixth Workshop on Very Large Corpora*, pages 161–170.
- Kameyama, Megumi. 1997. Recognizing referential links: An information extraction perspective. In *Proceedings of the ACL Workshop on Operational Factors in Practical, Robust Anaphora Resolution for Unrestricted Texts*, pages 46–53.
- Kehler, Andrew. 1997. Probabilistic coreference in information extraction. In *Proceedings of the Second Conference on Empirical Methods in Natural Language Processing*, pages 163–173.
- Lappin, Shalom and Herbert J. Leass. 1994. An algorithm for pronominal anaphora resolution. *Computational Linguistics*, 20(4):535–561, December.
- McCarthy, Joseph F. 1996. *A Trainable Approach to Coreference Resolution for Information Extraction*. Ph.D. thesis,

- University of Massachusetts Amherst,
Department of Computer Science,
September.
- McCarthy, Joseph F. and Wendy Lehnert.
1995. Using decision trees for coreference
resolution. In *Proceedings of the Fourteenth
International Joint Conference on Artificial
Intelligence*, pages 1050–1055.
- McEnery, A., I. Tanaka, and S. Botley. 1997.
Corpus annotation and reference
resolution. In *Proceedings of the ACL
Workshop on Operational Factors in Practical,
Robust Anaphora Resolution for Unrestricted
Texts*, pages 67–74.
- Miller, George A. 1990. WordNet: An
on-line lexical database. *International
Journal of Lexicography*, 3(4):235–312.
- Mitkov, Ruslan. 1997. Factors in anaphora
resolution: They are not the only things
that matter. A case study based on two
different approaches. In *Proceedings of the
ACL Workshop on Operational Factors in
Practical, Robust Anaphora Resolution for
Unrestricted Texts*, pages 14–21.
- MUC-6. 1995. Coreference task definition
(v2.3, 8 Sep 95). In *Proceedings of the Sixth
Message Understanding Conference (MUC-6)*,
pages 335–344.
- MUC-7. 1997. Coreference task definition
(v3.0, 13 Jul 97). In *Proceedings of the
Seventh Message Understanding Conference
(MUC-7)*.
- Quinlan, John Ross. 1993. *C4.5: Programs for
Machine Learning*. Morgan Kaufmann, San
Francisco, CA.
- Sundheim, Beth M. 1995. Overview of
results of the MUC-6 evaluation. In
*Proceedings of the Sixth Message
Understanding Conference (MUC-6)*, pages
13–31.

This article has been cited by:

1. Shasha Deng, Yilu Zhou, Pengzhu Zhang, Ahmed Abbasi. 2018. Using discussion logic in analyzing online group discussions: A text mining approach. *Information & Management* . [[Crossref](#)]
2. Manel Achichi, Mohamed Ben Ellefi, Zohra Bellahsene, Konstantin Todorov. Doing Web Data: from Dataset Recommendation to Data Linking 57-91. [[Crossref](#)]
3. Tugba Pamay, Gulsen Eryigit. Turkish Coreference Resolution 1-7. [[Crossref](#)]
4. Yufang Hou, Katja Markert, Michael Strube. 2018. Unrestricted Bridging Resolution. *Computational Linguistics* **44**:2, 237-284. [[Abstract](#)] [[Full Text](#)] [[PDF](#)] [[PDF Plus](#)]
5. Yufang Hou, Katja Markert, Michael Strube. Unrestricted Bridging Resolution. *Computational Linguistics*, ahead of print 237-284. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
6. David Chiang, Frank Drewes, Daniel Gildea, Adam Lopez, Giorgio Satta. 2018. Weighted DAG Automata for Semantic Graphs. *Computational Linguistics* **44**:1, 119-186. [[Abstract](#)] [[Full Text](#)] [[PDF](#)] [[PDF Plus](#)]
7. Sheng Cheng, Kong Fang, Zhou Guodong. Towards Better Chinese Zero Pronoun Resolution from Discourse Perspective 406-418. [[Crossref](#)]
8. Natalia Garanina, Elena Sidorova, Irina Kononenko. A Distributed Approach to Coreference Resolution in Multiagent Text Analysis for Ontology Population 147-162. [[Crossref](#)]
9. Gongqing Wu, Ying He, Xuegang Hu. 2018. Entity Linking: An Issue to Extract Corresponding Entity With Knowledge Base. *IEEE Access* **6**, 6220-6231. [[Crossref](#)]
10. Adèle Désoyer, Frédéric Landragin, Isabelle Tellier, Anaïs Lefeuvre, Jean-Yves Antoine, Marco Dinarelli. Coreference Resolution for French Oral Data: Machine Learning Experiments with ANCOR 507-519. [[Crossref](#)]
11. Souha Mezghani Hammami, Lamia Hadrich Belguith. Arabic Pronominal Anaphora Resolution Based on New Set of Features 533-544. [[Crossref](#)]
12. N. O. Garanina, E. A. Sidorova, A. S. Seryi. 2018. Multiagent Approach to Coreference Resolution Based on the Multifactor Similarity in Ontology Population. *Programming and Computer Software* **44**:1, 23-34. [[Crossref](#)]
13. Evandro Fonseca, Aline Vanin, Renata Vieira. Mention Clustering to Improve Portuguese Semantic Coreference Resolution 256-263. [[Crossref](#)]
14. Evandro Fonseca, Aline Vanin, Renata Vieira. Nominal Coreference Resolution Using Semantic Knowledge 37-45. [[Crossref](#)]
15. Nouf Bindris, Saatviga Sudhahar, Nello Cristianini. Fact Checking from Natural Text with Probabilistic Soft Logic 52-61. [[Crossref](#)]
16. R. Vijay Sundar Ram, Sobha Lalitha Devi. A Robust Coreference Chain Builder for Tamil 233-246. [[Crossref](#)]
17. Jose L. Martinez-Rodriguez, Aidan Hogan, Ivan Lopez-Arevalo, Andreas Hotho. 2018. Information extraction meets the Semantic Web: A survey. *Semantic Web* **1**. [[Crossref](#)]
18. Ajay Gupta, Devendra Verma, Sachin Pawar, Sangameshwar Patil, Swapnil Hingmire, Girish K. Palshikar, Pushpak Bhattacharyya. Identifying Participant Mentions and Resolving Their Coreferences in Legal Court Judgements 153-162. [[Crossref](#)]
19. Chen Li, Zhiqiang Rao, Qinghua Zheng, Xiangrong Zhang. 2018. A set of domain rules and a deep network for protein coreference resolution. *Database* **2018**. . [[Crossref](#)]
20. K. Bretonnel Cohen, Arrick Lanfranchi, Miji Joo-young Choi, Michael Bada, William A. Baumgartner, Natalya Panteleyeva, Karin Verspoor, Martha Palmer, Lawrence E. Hunter. 2017. Coreference annotation and resolution in the Colorado Richly Annotated Full Text (CRAFT) corpus of biomedical journal articles. *BMC Bioinformatics* **18**:1. . [[Crossref](#)]
21. . Bibliography 263-300. [[Crossref](#)]
22. Guimin Huang, Min Tan, Sirui Huang, Ruyu Mo, Ya Zhou. A discourse coherence model for analyzing Chinese students' essay 430-434. [[Crossref](#)]
23. Alexandra Khadzhiskaia, Andrey Sysoev. Coreference Resolution for Russian: Taking Stock and Moving Forward 70-75. [[Crossref](#)]
24. HEEYOUNG LEE, MIHAI SURDEANU, DAN JURAFSKY. 2017. A scaffolding approach to coreference resolution integrating statistical and rule-based models. *Natural Language Engineering* **23**:05, 733-762. [[Crossref](#)]
25. Tian Shengwei, Li Dongbai, Yu Long, Feng Guanjuan, Zhao Jianguo, Li Pu. 2017. Anaphoricity Determination of Anaphora Resolution in Uyghur Pronoun Based on CNN-LSTM Model. *International Journal of Computational Intelligence and Applications* **16**:03, 1750016. [[Crossref](#)]

26. G Veena, Deepa Gupta, Anna Neethu Daniel, S. Roshny. A learning method for coreference resolution using semantic role labeling features 67-72. [[Crossref](#)]
27. Changki Lee, Sangkeun Jung, Cheon-Eum Park. 2017. Anaphora resolution with pointer networks. *Pattern Recognition Letters* **95**, 1-7. [[Crossref](#)]
28. Paul Swoboda, Jan Kuske, Bogdan Savchynskyy. A Dual Ascent Framework for Lagrangean Decomposition of Combinatorial Problems 4950-4960. [[Crossref](#)]
29. Paul Swoboda, Bjoern Andres. A Message Passing Algorithm for the Minimum Cost Multicut Problem 4990-4999. [[Crossref](#)]
30. ANDER SORALUZE, OLATZ ARREGI, XABIER ARREGI, ARANTZA DÍAZ DE ILARRAZA. 2017. Improving mention detection for Basque based on a deep error analysis. *Natural Language Engineering* **23**:03, 351-384. [[Crossref](#)]
31. Bryan A. Plummer, Liwei Wang, Chris M. Cervantes, Juan C. Caicedo, Julia Hockenmaier, Svetlana Lazebnik. 2017. Flickr30k Entities: Collecting Region-to-Phrase Correspondences for Richer Image-to-Sentence Models. *International Journal of Computer Vision* **123**:1, 74-93. [[Crossref](#)]
32. Timothy Miller, Dmitriy Dligach, Steven Bethard, Chen Lin, Guergana Savova. 2017. Towards generalizable entity-centric clinical coreference resolution. *Journal of Biomedical Informatics* **69**, 251-258. [[Crossref](#)]
33. Lei Zou, M. Tamer Özsu. 2017. Graph-Based RDF Data Management. *Data Science and Engineering* **18**. . [[Crossref](#)]
34. Jennifer Lautenschlager, James Starz, Ian Warfield. A Statistical Approach to the Subnational Geolocation of Event Data 333-343. [[Crossref](#)]
35. Jheng-Long Wu, Wei-Yun Ma. A Deep Learning Framework for Coreference Resolution Based on Convolutional Neural Network 61-64. [[Crossref](#)]
36. Ryu Iida, Mamoru Komachi, Naoya Inoue, Kentaro Inui, Yuji Matsumoto. NAIST Text Corpus: Annotating Predicate- Argument and Coreference Relations in Japanese 1177-1196. [[Crossref](#)]
37. Gil Rocha, Henrique Lopes Cardoso. Towards a Mention-Pair Model for Coreference Resolution in Portuguese 855-867. [[Crossref](#)]
38. Thomas Lansdall-Welfare, Saatviga Sudhahar, James Thompson, Nello Cristianini. The Actors of History: Narrative Network Analysis Reveals the Institutions of Power in British Society Between 1800-1950 186-197. [[Crossref](#)]
39. Lili Jiang. Entity Markup for Knowledge Base Population 71-89. [[Crossref](#)]
40. Aria Nourbakhsh, Mohammad Bahrani. Persian Pronoun Resolution Using Data Driven Approaches 574-585. [[Crossref](#)]
41. Halil Kilicoglu, Graciela Rosemblat, Marcelo Fiszman, Thomas C. Rindflesch. 2016. Sortal anaphora resolution to enhance relation extraction from biomedical literature. *BMC Bioinformatics* **17**:1. . [[Crossref](#)]
42. Richa Sharma, Nidhi Sharma, K.K. Biswas. Machine Learning for Detecting Pronominal Anaphora Ambiguity in NL Requirements 177-182. [[Crossref](#)]
43. Sandra Kübler, Desislava Zhekova. 2016. Multilingual coreference resolution. *Language and Linguistics Compass* **10**:11, 614-631. [[Crossref](#)]
44. Sergio Oramas, Luis Espinosa-Anke, Mohamed Sordo, Horacio Saggion, Xavier Serra. 2016. Information extraction for knowledge base construction in the music domain. *Data & Knowledge Engineering* **106**, 70-83. [[Crossref](#)]
45. Utpal Kumar Sikdar, Asif Ekbal, Sriparna Saha. 2016. A generalized framework for anaphora resolution in Indian languages. *Knowledge-Based Systems* **109**, 147-159. [[Crossref](#)]
46. Anastasiya Bodrova, Natalia Grafeeva. Coreference resolution using clusterization 1-8. [[Crossref](#)]
47. Dipankar Das, Sunit Bhattacharya. Coreference between subjective expression and holder: A classification perspective 535-541. [[Crossref](#)]
48. Kenneth Joseph, Wei Wei, Matthew Benigni, Kathleen M. Carley. 2016. A social-event based approach to sentiment analysis of identities and behaviors in text. *The Journal of Mathematical Sociology* **40**:3, 137-166. [[Crossref](#)]
49. Miji Choi, Justin Zobel, Karin Verspoor. 2016. A categorical analysis of coreference resolution errors in biomedical texts. *Journal of Biomedical Informatics* **60**, 309-318. [[Crossref](#)]
50. Halil Kilicoglu, Dina Demner-Fushman. 2016. Bio-SCoRes: A Smorgasbord Architecture for Coreference Resolution in Biomedical Text. *PLOS ONE* **11**:3, e0148538. [[Crossref](#)]
51. Erica Briscoe, Scott Appling, Joel Schlosser. 2016. Technology Futures From Passive Crowdsourcing. *IEEE Transactions on Computational Social Systems* **3**:1, 23-31. [[Crossref](#)]
52. Piyoros Tungthamthiti, Kiyoaki Shirai, Masnizah Mohd. 2016. Recognition of Sarcasm in Microblogging Based on Sentiment Analysis and Coherence Identification. *Journal of Natural Language Processing* **23**:5, 383-405. [[Crossref](#)]

53. Veronique Hoste. The Mention-Pair Model 269-282. [[Crossref](#)]
54. Evandro Fonseca, Renata Vieira, Aline Vanin. Improving Coreference Resolution with Semantic Knowledge 213-224. [[Crossref](#)]
55. Bingbing Wu, Tiejun Zhao. A Context-Aware Model Using Distributed Representations for Chinese Zero Pronoun Resolution 3-11. [[Crossref](#)]
56. Yannick Versley, Massimo Poesio, Simone Ponzetto. Using Lexical and Encyclopedic Knowledge 393-429. [[Crossref](#)]
57. Massimo Poesio, Roland Stuckardt, Yannick Versley, Renata Vieira. Early Approaches to Anaphora Resolution: Theoretically Inspired and Heuristic-Based 55-94. [[Crossref](#)]
58. Marta Recasens, Sameer Pradhan. Evaluation Campaigns 165-208. [[Crossref](#)]
59. Nick Rizzolo, Dan Roth. Integer Linear Programming for Coreference Resolution 315-343. [[Crossref](#)]
60. Shane Bergsma. Extracting Anaphoric Agreement Properties from Corpora 345-368. [[Crossref](#)]
61. Vincent Ng. Advanced Machine Learning Models for Coreference Resolution 283-313. [[Crossref](#)]
62. Xianmin Liu, Jianzhong Li, Hong Gao. 2016. On the hardness of labeled correlation clustering problem: A parameterized complexity view. *Theoretical Computer Science* **609**, 583-593. [[Crossref](#)]
63. Thi Thuy Le, Thanh Hung Vo, Duc Trung Mai, Than Tho Quan, Tuoi Thi Phan. Sentiment Analysis Using Anaphoric Coreference Resolution and Ontology Inference 297-303. [[Crossref](#)]
64. Lavika Goel, Rashi Khandelwal, Eloy Retamino, Suraj Nair, Alois Knoll. Application of a Hybrid Relation Extraction Framework for Intelligent Natural Language Processing 803-813. [[Crossref](#)]
65. Yannick Versley, Anders Björkelund. Off-the-Shelf Tools 237-266. [[Crossref](#)]
66. Roland Stuckardt. Introduction 1-19. [[Crossref](#)]
67. Xue-feng Xi, Guodong Zhou. A Micro-topic Model for Coreference Resolution Based on Theme-Rheme Structure 648-656. [[Crossref](#)]
68. Laure Berti-Équille, Javier Borge-Holthoefer. 2015. Veracity of Data: From Truth Discovery Computation Algorithms to Models of Misinformation Dynamics. *Synthesis Lectures on Data Management* **7:3**, 1-155. [[Crossref](#)]
69. Bryan A. Plummer, Liwei Wang, Chris M. Cervantes, Juan C. Caicedo, Julia Hockenmaier, Svetlana Lazebnik. Flickr30k Entities: Collecting Region-to-Phrase Correspondences for Richer Image-to-Sentence Models 2641-2649. [[Crossref](#)]
70. Mayank Kejriwal, Daniel P. Miranker. 2015. An unsupervised instance matcher for schema-free RDF data. *Journal of Web Semantics* **35**, 102-123. [[Crossref](#)]
71. Sebastian Martschat, Michael Strube. 2015. Latent Structures for Coreference Resolution. *Transactions of the Association for Computational Linguistics* **3**, 405-418. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
72. Bernardo Cabaleiro, Anselmo Peñas. 2015. On improving parsing with automatically acquired semantic classes. *Knowledge-Based Systems* **89**, 359-365. [[Crossref](#)]
73. Ana Zelaia, Olatz Arregi, Basilio Sierra. 2015. Combining Singular Value Decomposition and a multi-classifier: A new approach to support coreference resolution. *Engineering Applications of Artificial Intelligence* **46**, 279-286. [[Crossref](#)]
74. Qiang Li, Zongtian Liu, Lei Chen, Xianchuan Wang. An Event-Oriented Multi-pass Sieve Module for Coreference Resolution 203-207. [[Crossref](#)]
75. Utpal Kumar Sikdar, Asif Ekbal, Sriparna Saha, Olga Uryupina, Massimo Poesio. 2015. Differential evolution-based feature selection technique for anaphora resolution. *Soft Computing* **19:8**, 2149-2161. [[Crossref](#)]
76. GORAN GLAVAŠ, JAN ŠNAJDER. 2015. Construction and evaluation of event graphs. *Natural Language Engineering* **21:04**, 607-652. [[Crossref](#)]
77. Min-Kook Hwang, Youngtae Kim, Dongyul Ra, Soojong Lim, Hyunki Kim. 2015. Restoring Omitted Sentence Constituents in Encyclopedia Documents Using Structural SVM. *Journal of Intelligence and Information Systems* **21:2**, 131-150. [[Crossref](#)]
78. Zhao-Yan Ming, Tat Seng Chua. 2015. Resolving polysemy and pseudonymity in entity linking with comprehensive name and context modeling. *Information Sciences* **307**, 18-38. [[Crossref](#)]
79. Yanyan Zhao, Bing Qin, Ting Liu, Wei Yang. 2015. Aspect-Object Alignment with Integer Linear Programming in Opinion Mining. *PLOS ONE* **10:5**, e0125084. [[Crossref](#)]
80. Slavko Zitnik, Marko Bajec. Iterative joint extraction of entities, relationships and coreferences from text sources 412-422. [[Crossref](#)]
81. John Atkinson, Gonzalo Salas, Alejandro Figueroa. 2015. Improving opinion retrieval in social media by combining features-based coreferencing and memory-based learning. *Information Sciences* **299**, 20-31. [[Crossref](#)]

82. Xianmin Liu, Jianzhong Li. 2015. Algorithms and complexity results for labeled correlation clustering problem. *Journal of Combinatorial Optimization* **29**:2, 488-501. [[Crossref](#)]
83. Xiaodan Zhu, Peter Turney, Daniel Lemire, André Vellino. 2015. Measuring academic influence: Not all citations are equal. *Journal of the Association for Information Science and Technology* **66**:2, 408-427. [[Crossref](#)]
84. SAATVIGA SUDHAHAR, GIANLUCA DE FAZIO, ROBERTO FRANZOSI, NELLO CRISTIANINI. 2015. Network analysis of narrative content in large corpora. *Natural Language Engineering* **21**:01, 81-112. [[Crossref](#)]
85. Sibow Wang, Xiaokui Xiao, Chun-Hee Lee. Crowd-Based Deduplication 1263-1277. [[Crossref](#)]
86. Roman Prokofyev, Alberto Tonon, Michael Luggen, Loic Vouilloz, Djellel Eddine Difallah, Philippe Cudré-Mauroux. SANAPHOR: Ontology-Based Coreference Resolution 458-473. [[Crossref](#)]
87. Xue-Feng Xi, Guodong Zhou, Fuyuan Hu, Bao-chuan Fu. A Convolutional Deep Neural Network for Coreference Resolution via Modeling Hierarchical Features 361-372. [[Crossref](#)]
88. Voldemaras Žitkus, Lina Nemuraitė. First Steps in Automatic Anaphora Resolution in Lithuanian Language Based on Morphological Annotations and Named Entity Recognition 480-490. [[Crossref](#)]
89. Eraldo Rezende Fernandes, Cícero Nogueira dos Santos, Ruy Luiz Milidiú. 2014. Latent Trees for Coreference Resolution. *Computational Linguistics* **40**:4, 801-835. [[Abstract](#)] [[Full Text](#)] [[PDF](#)] [[PDF Plus](#)]
90. Paul Felt, Eric K. Ringger, Kevin Seppi, Kristian S. Heal, Robbie A. Haertel, Deryle Lonsdale. 2014. Evaluating machine-assisted annotation in under-resourced settings. *Language Resources and Evaluation* **48**:4, 561-599. [[Crossref](#)]
91. Greg Durrett, Dan Klein. 2014. A Joint Model for Entity Analysis: Coreference, Typing, and Linking. *Transactions of the Association for Computational Linguistics* **2**, 477-490. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
92. Ankur Agarwal, Ravi S. Behara, Sirish Mulpura, Vivek Tyagi. Domain Independent Natural Language Processing -- A Case Study for Hospital Readmission with COPD 399-404. [[Crossref](#)]
93. Chih-Hao Ku, GONDY Leroy. 2014. A decision support system: Automated crime report analysis and classification for e-government. *Government Information Quarterly* **31**:4, 534-544. [[Crossref](#)]
94. Sharon Gower Small, Larry Medsker. 2014. Review of information extraction technologies and applications. *Neural Computing and Applications* **25**:3-4, 533-548. [[Crossref](#)]
95. Slavko Žitnik, Lovro Šubelj, Marko Bajec. 2014. SkipCor: Skip-Mention Coreference Resolution Using Linear-Chain Conditional Random Fields. *PLoS ONE* **9**:6, e100101. [[Crossref](#)]
96. Oanh Thi Tran, Bach Xuan Ngo, Minh Le Nguyen, Akira Shimazu. 2014. Automated reference resolution in legal texts. *Artificial Intelligence and Law* **22**:1, 29-60. [[Crossref](#)]
97. Detlef Schoder, Johannes Putzke, Panagiotis Takis Metaxas, Peter A. Gloor, Kai Fischbach. 2014. Information Systems for "Wicked Problems". *Business & Information Systems Engineering* **6**:1, 3-10. [[Crossref](#)]
98. Detlef Schoder, Johannes Putzke, Panagiotis Takis Metaxas, Peter A. Gloor, Kai Fischbach. 2014. Informationssysteme für „Wicked Problems“. *WIRTSCHAFTSINFORMATIK* **56**:1, 3-11. [[Crossref](#)]
99. Imed Zitouni, Yassine Benajiba. 2014. Aligned-Parallel-Corpora Based Semi-Supervised Learning for Arabic Mention Detection. *IEEE/ACM Transactions on Audio, Speech, and Language Processing* **22**:2, 314-324. [[Crossref](#)]
100. Xianmin Liu, Jianzhong Li, Hong Gao. On the Parameterized Complexity of Labelled Correlation Clustering Problem 276-287. [[Crossref](#)]
101. Jong C. Park, Hee-Jin Lee. Augmenting Biological Text Mining with Symbolic Inference 901-918. [[Crossref](#)]
102. Emili Sapena, Lluís Padró, Jordi Turmo. 2013. A Constraint-Based Hypergraph Partitioning Approach to Coreference Resolution. *Computational Linguistics* **39**:4, 847-884. [[Abstract](#)] [[Full Text](#)] [[PDF](#)] [[PDF Plus](#)]
103. Heeyoung Lee, Angel Chang, Yves Peirsman, Nathanael Chambers, Mihai Surdeanu, Dan Jurafsky. 2013. Deterministic Coreference Resolution Based on Entity-Centric, Precision-Ranked Rules. *Computational Linguistics* **39**:4, 885-916. [[Abstract](#)] [[Full Text](#)] [[PDF](#)] [[PDF Plus](#)]
104. Hong-Jie Dai, Johnny Chi-Yang Wu, Richard Tzong-Han Tsai. 2013. Collective Instance-Level Gene Normalization on the IGN Corpus. *PLoS ONE* **8**:11, e79517. [[Crossref](#)]
105. Kristina Nilsson Björkenstam. 2013. SUC-CORE: A Balanced Corpus Annotated with Noun Phrase Coreference. *Northern European Journal of Language Technology* **3**, 19-39. [[Crossref](#)]
106. Lluís Màrquez, Marta Recasens, Emili Sapena. 2013. Coreference resolution: an empirical study based on SemEval-2010 shared Task 1. *Language Resources and Evaluation* **47**:3, 661-694. [[Crossref](#)]

107. Amit Dubey, Frank Keller, Patrick Sturt. 2013. Probabilistic Modeling of Discourse-Aware Sentence Processing. *Topics in Cognitive Science* 5:3, 425-451. [[Crossref](#)]
108. Robert Frank, Donald Mathis, William Badecker. 2013. The Acquisition of Anaphora by Simple Recurrent Networks. *Language Acquisition* 20:3, 181-227. [[Crossref](#)]
109. Md. Faisal Mahbub Chowdhury, Pierre Zweigenbaum. 2013. A controlled greedy supervised approach for co-reference resolution on clinical text. *Journal of Biomedical Informatics* 46:3, 506-515. [[Crossref](#)]
110. C. Cig, T. M. Sezgin. New modalities, new challenges - Annotating sketching and gaze data 1-4. [[Crossref](#)]
111. K. Karthikeyan, V. Karthikeyani. Understanding text using Anaphora Resolution 346-350. [[Crossref](#)]
112. Xi Peng, Zhang Yi, Xiao-Yong Wei, De-Zhong Peng, Yong-Sheng Sang. 2013. Free-gram phrase identification for modeling Chinese text. *Information Processing Letters* 113:4, 137-144. [[Crossref](#)]
113. David Mendes, Irene Pimenta Rodrigues. A Semantic Web Pragmatic Approach to Develop Clinical Ontologies, and thus Semantic Interoperability, based in HL7 v2.xml Messaging 205-214. [[Crossref](#)]
114. Iris Hendrickx, Gosse Bouma, Walter Daelemans, Véronique Hoste. COREA: Coreference Resolution for Extracting Answers for Dutch 115-128. [[Crossref](#)]
115. Ngan Nguyen, Jin-Dong Kim, Makoto Miwa, Takuya Matsuzaki, Junichi Tsujii. 2012. Improving protein coreference resolution by simple semantic classification. *BMC Bioinformatics* 13:1. . [[Crossref](#)]
116. Yan Xu, Jiahua Liu, Jiajun Wu, Yue Wang, Zhuowen Tu, Jian-Tao Sun, Junichi Tsujii, Eric I-Chao Chang. 2012. A classification approach to coreference in discharge summaries: 2011 i2b2 challenge. *Journal of the American Medical Informatics Association* 19:5, 897-905. [[Crossref](#)]
117. Andreea Bodnari, Peter Szolovits, Özlem Uzuner. 2012. MCORES: a system for noun phrase coreference resolution for clinical records. *Journal of the American Medical Informatics Association* 19:5, 906-912. [[Crossref](#)]
118. Siddhartha Reddy Jonnalagadda, Dingcheng Li, Sunghwan Sohn, Stephen Tze-Inn Wu, Kavishwar Waghlikar, Manabu Torii, Hongfang Liu. 2012. Coreference analysis in clinical notes: a multi-pass sieve with alternate anaphora resolution modules. *Journal of the American Medical Informatics Association* 19:5, 867-874. [[Crossref](#)]
119. Hong-Jie Dai, Chun-Yu Chen, Chi-Yang Wu, Po-Ting Lai, Richard Tzong-Han Tsai, Wen-Lian Hsu. 2012. Coreference resolution of medical concepts in discharge summaries by exploiting contextual information. *Journal of the American Medical Informatics Association* 19:5, 888-896. [[Crossref](#)]
120. Helene Schmolz, David Coquil, Mario Doller. In-Depth Analysis of Anaphora Resolution Requirements 174-179. [[Crossref](#)]
121. Elena Zheleva, Evimaria Terzi, Lise Getoor. 2012. Privacy in Social Networks. *Synthesis Lectures on Data Mining and Knowledge Discovery* 3:1, 1-85. [[Crossref](#)]
122. Shuhan Qi, Xuan Wang, Xinxin Li. Improve Coreference Resolution with Parameter Tunable Anaphoricity Identification and Global Optimization 307-314. [[Crossref](#)]
123. R. Vijay Sundar Ram, Sobha Lalitha Devi. Coreference Resolution Using Tree CRFs 285-296. [[Crossref](#)]
124. Jiaping Zheng, Wendy W. Chapman, Rebecca S. Crowley, Guergana K. Savova. 2011. Coreference resolution: A review of general methodologies and applications in the clinical domain. *Journal of Biomedical Informatics* 44:6, 1113-1122. [[Crossref](#)]
125. Manfred Stede. 2011. Discourse Processing. *Synthesis Lectures on Human Language Technologies* 4:3, 1-165. [[Crossref](#)]
126. Junwei Gao, Fang Kong, Peifeng Li, Qiaoming Zhu. Research of Noun Phrase Coreference Resolution 93-96. [[Crossref](#)]
127. Ning Zhang, Fang Kong, Peifeng Li. Research of Event Pronoun Resolution 85-88. [[Crossref](#)]
128. M. RECASENS, E. HOVY. 2011. BLANC: Implementing the Rand index for coreference evaluation. *Natural Language Engineering* 17:04, 485-510. [[Crossref](#)]
129. Hong-Jie Dai, Yen-Ching Chang, Richard Tzong-Han Tsai, Wen-Lian Hsu. 2011. Integration of gene normalization stages and co-reference resolution using a Markov logic network. *Bioinformatics* 27:18, 2586-2594. [[Crossref](#)]
130. Hui Yang, Anne de Roeck, Vincenzo Gervasi, Alistair Willis, Bashar Nuseibeh. 2011. Analysing anaphoric ambiguity in natural language requirements. *Requirements Engineering* 16:3, 163-189. [[Crossref](#)]
131. Zhong-Chao FEI, Ya-Qian ZHOU, Xuan-Jing HUANG, Li-De WU. 2011. Pronoun Resolution in Spoken Dialog. *Journal of Software* 22:2, 233-244. [[Crossref](#)]
132. Kaihong Liu, William R. Hogan, Rebecca S. Crowley. 2011. Natural Language Processing methods and systems for biomedical ontology learning. *Journal of Biomedical Informatics* 44:1, 163-179. [[Crossref](#)]

133. Yang Yang, Priyanka Singh, Jiadi Yao, Ching-man Au Yeung, Amir Zareian, Xiaowei Wang, Zhonglun Cai, Manuel Salvadorés, Nicholas Gibbins, Wendy Hall, Nigel Shadbolt. Distributed Human Computation Framework for Linked Data Co-reference Resolution 32-46. [[Crossref](#)]
134. Guo-Dong Zhou, Fang Kong. 2011. Learning Noun Phrase Anaphoricity in Coreference Resolution via Label Propagation. *Journal of Computer Science and Technology* **26**:1, 34-44. [[Crossref](#)]
135. Katsumasa Yoshikawa, Tsutomu Hirao, Sebastian Riedel, Masayuki Asahara, Yuji Matsumoto. 2011. Coreference Based Event Extraction on Biomedical Text. *Transactions of the Japanese Society for Artificial Intelligence* **26**, 318-323. [[Crossref](#)]
136. Ben Hachey, Will Radford, James R. Curran. Graph-Based Named Entity Linking with Wikipedia 213-226. [[Crossref](#)]
137. Michal Novák, Zdeněk Žabokrtský. Resolving Noun Phrase Coreference in Czech 24-34. [[Crossref](#)]
138. José Guilherme Camargo de Souza, Constantin Orăsan. Can Projected Chains in Parallel Corpora Help Coreference Resolution? 59-69. [[Crossref](#)]
139. Iris Hendrickx, Orphée De Clercq, Veronique Hoste. Analysis and Reference Resolution of Bridge Anaphora across Different Text Genres 1-11. [[Crossref](#)]
140. Katsumasa Yoshikawa, Sebastian Riedel, Tsutomu Hirao, Masayuki Asahara, Yuji Matsumoto. 2011. Coreference based event-argument relation extraction on biomedical text. *Journal of Biomedical Semantics* **2**:Suppl 5, S6. [[Crossref](#)]
141. David Mendes, Irene Rodrigues. A Semantic Web Pragmatic Approach to Develop Clinical Ontologies, and Thus Semantic Interoperability, Based in HL7 v2.XML Messaging 297-306. [[Crossref](#)]
142. Asif Ekbal, Sriparna Saha, Olga Uryupina, Massimo Poesio. Multiobjective Simulated Annealing Based Approach for Feature Selection in Anaphora Resolution 47-58. [[Crossref](#)]
143. Amanda J. Stent, Srinivas Bangalore. Interaction between dialog structure and coreference resolution 342-347. [[Crossref](#)]
144. Marta Recasens, M. Antònia Martí. 2010. AnCora-CO: Coreferentially annotated corpora for Spanish and Catalan. *Language Resources and Evaluation* **44**:4, 315-345. [[Crossref](#)]
145. Fang Kong, Jianmei Zhou, Guodong Zhou, Qiaoming Zhu. Dependency Tree-Based Anaphoricity Determination for Coreference Resolution 215-218. [[Crossref](#)]
146. Cuili Huang, Yaqiang Wang, Yongmei Zhang, Yu Jin, Zhonghua Yu. Coreference resolution in biomedical full-text articles with domain dependent features 616-620. [[Crossref](#)]
147. Ratna Sanyal, Kushal Keshri, Vidya Nand. 2010. Importance of retrieving noun phrases and named entities from digital library content. *Journal of Zhejiang University SCIENCE C* **11**:11, 844-849. [[Crossref](#)]
148. Micha Elsner, Eugene Charniak. 2010. Disentangling Chat. *Computational Linguistics* **36**:3, 389-409. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
149. W. D. Lewis, F. Xia. 2010. Developing ODIN: A Multilingual Repository of Annotated Language Data for Hundreds of the World's Languages. *Literary and Linguistic Computing* **25**:3, 303-319. [[Crossref](#)]
150. Dong Jie. A Re-ranking Algorithm Based on Focused Named Entities 187-191. [[Crossref](#)]
151. . References 655-741. [[Crossref](#)]
152. Liang Zong, Xiaojun Wan, Lihong Zhao, Jianwu Yang, Yuqian Wu. Named Entity Resolution in Chinese News Comments on the Web 307-313. [[Crossref](#)]
153. Cheng-Yu Lu, Shian-Hua Lin, Jen-Chang Liu, Samuel Cruz-Lara, Jen-Shin Hong. 2010. Automatic event-level textual emotion sensing using mutual action histogram between entities. *Expert Systems with Applications* **37**:2, 1643-1653. [[Crossref](#)]
154. Volha Bryl, Claudio Giuliano, Luciano Serafini, Kateryna Tymoshenko. Supporting Natural Language Processing with Background Knowledge: Coreference Resolution Case 80-95. [[Crossref](#)]
155. Naoya Inoue, Ryu Iida, Kentaro Inui, Yuji Matsumoto. 2010. Resolving Direct and Indirect Anaphora for Japanese Definite Noun Phrases. *Journal of Natural Language Processing* **17**:1, 221-246. [[Crossref](#)]
156. Ryu Iida, Mamoru Komachi, Naoya Inoue, Kentaro Inui, Yuji Matsumoto. 2010. Annotating Predicate-Argument Relations and Anaphoric Relations: Findings from the Building of the NAIST Text Corpus. *Journal of Natural Language Processing* **17**:2, 25-50. [[Crossref](#)]
157. Ana Zelaia Jauregi, Basilio Sierra, Olatz Arregi Uriarte, Klara Ceberio, Arantza Díaz de Illaraza, Iakes Goenaga. A Combination of Classifiers for the Pronominal Anaphora Resolution in Basque 253-260. [[Crossref](#)]
158. O. Arregi, K. Ceberio, A. Díaz de Illaraza, I. Goenaga, B. Sierra, A. Zelaia. A First Machine Learning Approach to Pronominal Anaphora Resolution in Basque 234-243. [[Crossref](#)]

159. Yun-Soo Choi, Chang-Hoo Cheong, Sung-Pil Choi, Beom-Jong You, Jae-Hoon Kim. 2009. A Study on the Integration of Information Extraction Technology for Detecting Scientific Core Entities based on Large Resources. *Journal of Information Management* **40**:4, 1-22. [[Crossref](#)]
160. Fang Kong, Yancui Li, Guodong Zhou, Qiaoming Zhu. Exploring Syntactic Features for Pronoun Resolution Using Context-Sensitive Convolution Tree Kernel 201-205. [[Crossref](#)]
161. Yabing Zhang, Junsheng Zhou, Shujian Huang, Jiajun Chen. Combining ILP and MLN for Coreference Resolution 59-64. [[Crossref](#)]
162. Olena Medelyan, David Milne, Catherine Legg, Ian H. Witten. 2009. Mining meaning from Wikipedia. *International Journal of Human-Computer Studies* **67**:9, 716-754. [[Crossref](#)]
163. Imed Zitouni, Xiaoqiang Luo, Radu Florian. 2009. A Cascaded Approach to Mention Detection and Chaining in Arabic. *IEEE Transactions on Audio, Speech, and Language Processing* **17**:5, 935-944. [[Crossref](#)]
164. Yılmaz Kılıçaslan, Edip Serdar Güner, Savaş Yıldırım. 2009. Learning-based pronoun resolution for Turkish with a comparative evaluation. *Computer Speech & Language* **23**:3, 311-331. [[Crossref](#)]
165. Jun LANG. 2009. Number Type Recognition of Chinese Personal Noun Phrase. *Acta Automatica Sinica* **34**:8, 972-979. [[Crossref](#)]
166. Savaş Yıldırım, Yılmaz Kılıçaslan, Tuğba Yıldız. Pronoun Resolution in Turkish Using Decision Tree and Rule-Based Learning Algorithms 270-278. [[Crossref](#)]
167. Yan Chen, Yan-Qing Zhang. A Query Substitution-Search Result Refinement Approach for Long Query Web Searches 245-251. [[Crossref](#)]
168. Xiaofeng Yang, Jian Su, Chew Lim Tan. 2008. A Twin-Candidate Model for Learning-Based Anaphora Resolution. *Computational Linguistics* **34**:3, 327-356. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
169. Robert Elwell, Jason Baldridge. Discourse Connective Argument Identification with Connective Specific Rankers 198-205. [[Crossref](#)]
170. FEILIANG REN, JINGBO ZHU, HUIZHEN WANG, TONG XIAO. 2008. An Effective Approach for Coreference Resolution. *International Journal of Computer Processing of Languages* **21**:02, 135-149. [[Crossref](#)]
171. Tianjun Fu, Ahmed Abbasi, Hsinchun Chen. 2008. A hybrid approach to Web forum interactional coherence analysis. *Journal of the American Society for Information Science and Technology* **59**:8, 1195-1209. [[Crossref](#)]
172. Regina Barzilay, Mirella Lapata. 2008. Modeling Local Coherence: An Entity-Based Approach. *Computational Linguistics* **34**:1, 1-34. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
173. Fang Kong, Yancui Li, Guodong Zhou, Qiaoming Zhu, Peide Qian. Using Semantic Roles for Coreference Resolution 150-155. [[Crossref](#)]
174. Viktor Pekar. 2008. Discovery of event entailment knowledge from text corpora. *Computer Speech & Language* **22**:1, 1-16. [[Crossref](#)]
175. Ramon Ré Moya Cuevas, Willian Yukio Honda, Diego Jesus de Lucena, Ivandré Paraboni, Patrícia Rufino Oliveira. Portuguese Pronoun Resolution: Resources and Evaluation 344-350. [[Crossref](#)]
176. WENJIE LI, DONGLEI QIAN, QIN LU, CHUNFA YUAN. 2007. Chinese Entity Detection and Tracking: The Experience in ACE. *International Journal of Computer Processing of Languages* **20**:04, 219-236. [[Crossref](#)]
177. Sisay Fissaha Adafre, Valentin Jijkoun, Maarten de Rijke. Fact Discovery in Wikipedia 177-183. [[Crossref](#)]
178. Sameer S. Pradhan, Lance Ramshaw, Ralph Weischedel, Jessica MacBride, Linnea Micciulla. Unrestricted Coreference: Identifying Entities and Events in OntoNotes 446-453. [[Crossref](#)]
179. Lina Zhou. 2007. Ontology learning: state of the art and open issues. *Information Technology and Management* **8**:3, 241-252. [[Crossref](#)]
180. Jing Guan, Yanquan Zhou, Huacan He. A sort Approach for Anaphora Resolution of Chinese Personal Pronoun Based on Machine Learning Method 293-300. [[Crossref](#)]
181. Jun Zhao, Fei-Fan Liu. 2007. Linguistic Theory Based Contextual Evidence Mining for Statistical Chinese Co-Reference Resolution. *Journal of Computer Science and Technology* **22**:4, 608-617. [[Crossref](#)]
182. Ruslan Mitkov, Catalina Hallett. 2007. COMPARING PRONOUN RESOLUTION ALGORITHMS. *Computational Intelligence* **23**:2, 262-297. [[Crossref](#)]
183. Manabu Torii, K. Vijay-Shanker. 2007. SORTAL ANAPHORA RESOLUTION IN MEDLINE ABSTRACTS. *Computational Intelligence* **23**:1, 15-27. [[Crossref](#)]

184. Jason Baldridge, Nicholas Asher, Julie Hunter. 2007. Annotation for and Robust Parsing of Discourse Structure on Unrestricted Texts. *Zeitschrift für Sprachwissenschaft* 26:2. . [[Crossref](#)]
185. Donglei Qian, Wenjie Li, Chunfa Yuan, Qin Lu, Mingli Wu. Applying Machine Learning to Chinese Entity Detection and Tracking 154-165. [[Crossref](#)]
186. M. Strube. Anaphora and Coreference Resolution, Statistical 216-222. [[Crossref](#)]
187. Kwok Cheung Lan, Kei Shiu Ho, Robert Wing Pong Luk, Daniel So Yeung. 2005. FNDS: a dialogue-based system for accessing digested financial news. *Journal of Systems and Software* 78:2, 180-193. [[Crossref](#)]
188. Katja Markert, Malvina Nissim. 2005. Comparing Knowledge Sources for Nominal Anaphora Resolution. *Computational Linguistics* 31:3, 367-402. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
189. Bonnie Webber. 2004. D-LTAG: extending lexicalized TAG to discourse. *Cognitive Science* 28:5, 751-779. [[Crossref](#)]
190. Reviewed by Michael Strube. 2003. Book Review. *Computational Linguistics* 29:4, 662-664. [[Citation](#)] [[PDF](#)] [[PDF Plus](#)]
191. Derrick Higgins, Jerrold M. Sadock. 2003. A Machine Learning Approach to Modeling Scope Preferences. *Computational Linguistics* 29:1, 73-96. [[Abstract](#)] [[PDF](#)] [[PDF Plus](#)]
192. Zhiqiang Wang, Lei Li, Ruifan Li, Yixin Zhong. Chinese Pronominal Coreference Resolution Using Decision Tree Plus Filter Rules 587-591. [[Crossref](#)]
193. Ryu Iida, K. Inui, Y. Matsumoto. On the Issue of Combining Anaphoricity Determination and Antecedent Identification in Anaphora Resolution 244-249. [[Crossref](#)]
194. Jing Peng, Kenji Araki. Zero Anaphora Resolution in Chinese and Its Application in Chinese-English Machine Translation 364-375. [[Crossref](#)]
195. Nosheen Riaz, Moez Rehman. Negotiation by Software Agents in Electronic Business 327-349. [[Crossref](#)]
196. David Mendes, Irene Pimenta Rodrigues, Carlos Fernandes Baeta. Clinical Practice Ontology Automatic Learning from SOAP Reports 625-640. [[Crossref](#)]
197. Olga Uryupina, Massimo Poesio, Claudio Giuliano, Kateryna Tymoshenko. Disambiguation and Filtering Methods in Using Web Knowledge for Coreference Resolution 185-201. [[Crossref](#)]
198. David Mendes, Irene Pimenta Rodrigues, Carlos Fernandes Baeta. Clinical Practice Ontology Automatic Learning from SOAP Reports 349-363. [[Crossref](#)]
199. Qiang Shen, Tossapon Boongoen. Social Network Inspired Approach to Intelligent Monitoring of Intelligence Data 79-100. [[Crossref](#)]
200. Iris Hendrickx, Veronique Hoste, Walter Daelemans. Semantic and Syntactic Features for Dutch Coreference Resolution 351-361. [[Crossref](#)]